



High Performance LLM serving on NVIDIA GPUs

Sharan Chetlur, Principal Engineer, TRT-LLM | Event/Date

Production Language Apps

Deploying massive models for real-time applications

- Increasing need for deep learning in language applications
 - Chat, translation, summarization, search, generation, etc.
- **High accuracy** models are important for correct results
 - Model accuracy directly correlates to helpfulness for users
- “Online” deployment require **end-user acceptable latencies**
 - Ensure a great experience with applications
- Multi-functional, accurate models are *large* making them slow during inference & **expensive to deploy**

Making **cost effective** deployments **challenging**



Large Language Model Ecosystem

Rapid evolution makes optimization challenging

- Increasing rate of new foundational LLMs being released
 - Llama, Falcon, Starcoder, ChatGLM, MPT, & more
- New operators & customization techniques makes optimization a moving target
- Latest models continue to be very large for best accuracy
 - 70-200 Billion parameter or more

Need a **performant, robust, & extensible solution** for **cost-effective, real-time** LLM deployments

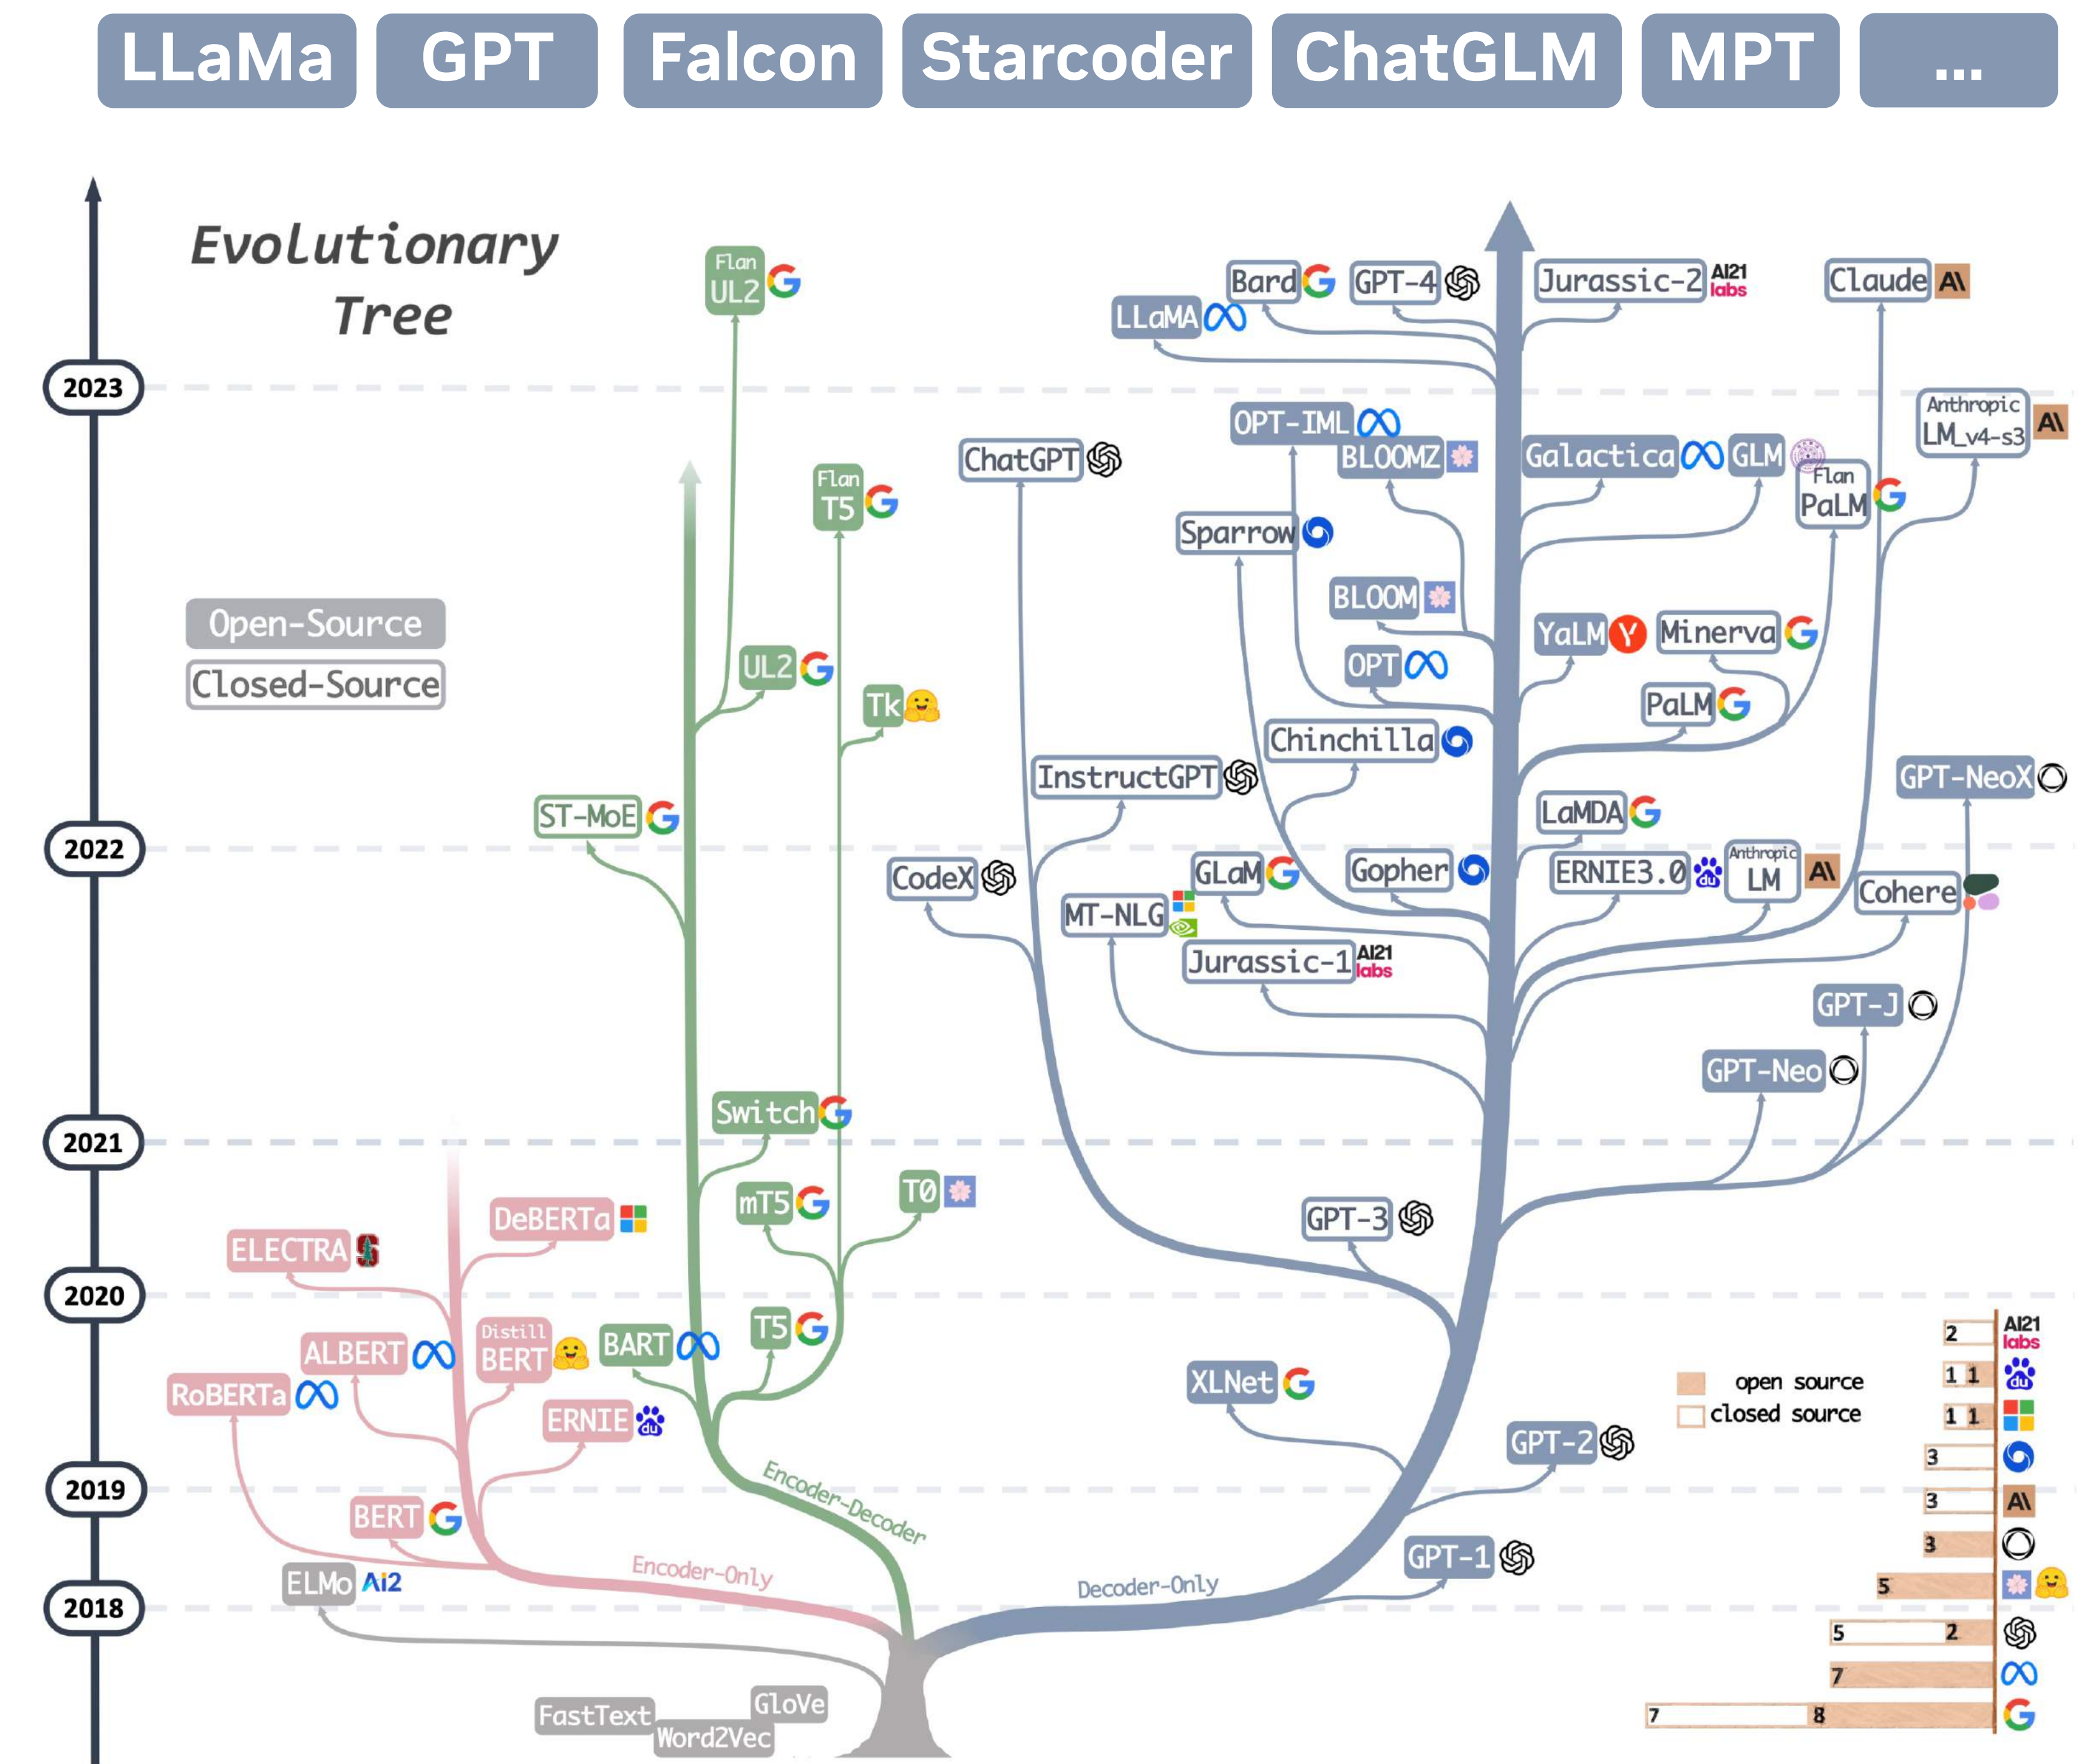


Image from [Mooler0410/LLMsPracticalGuide](https://mooler0410.github.io/LLMsPracticalGuide)

Yang, J., Jin, H., Tang, R., Han, X., Feng, Q., Jiang, H., ... Hu, X. (2023). Harnessing the Power of LLMs in Practice: A Survey on ChatGPT and Beyond. arXiv [Cs.CL]. Retrieved from <http://arxiv.org/abs/2304.13712>

TensorRT-LLM Optimizing LLM Inference

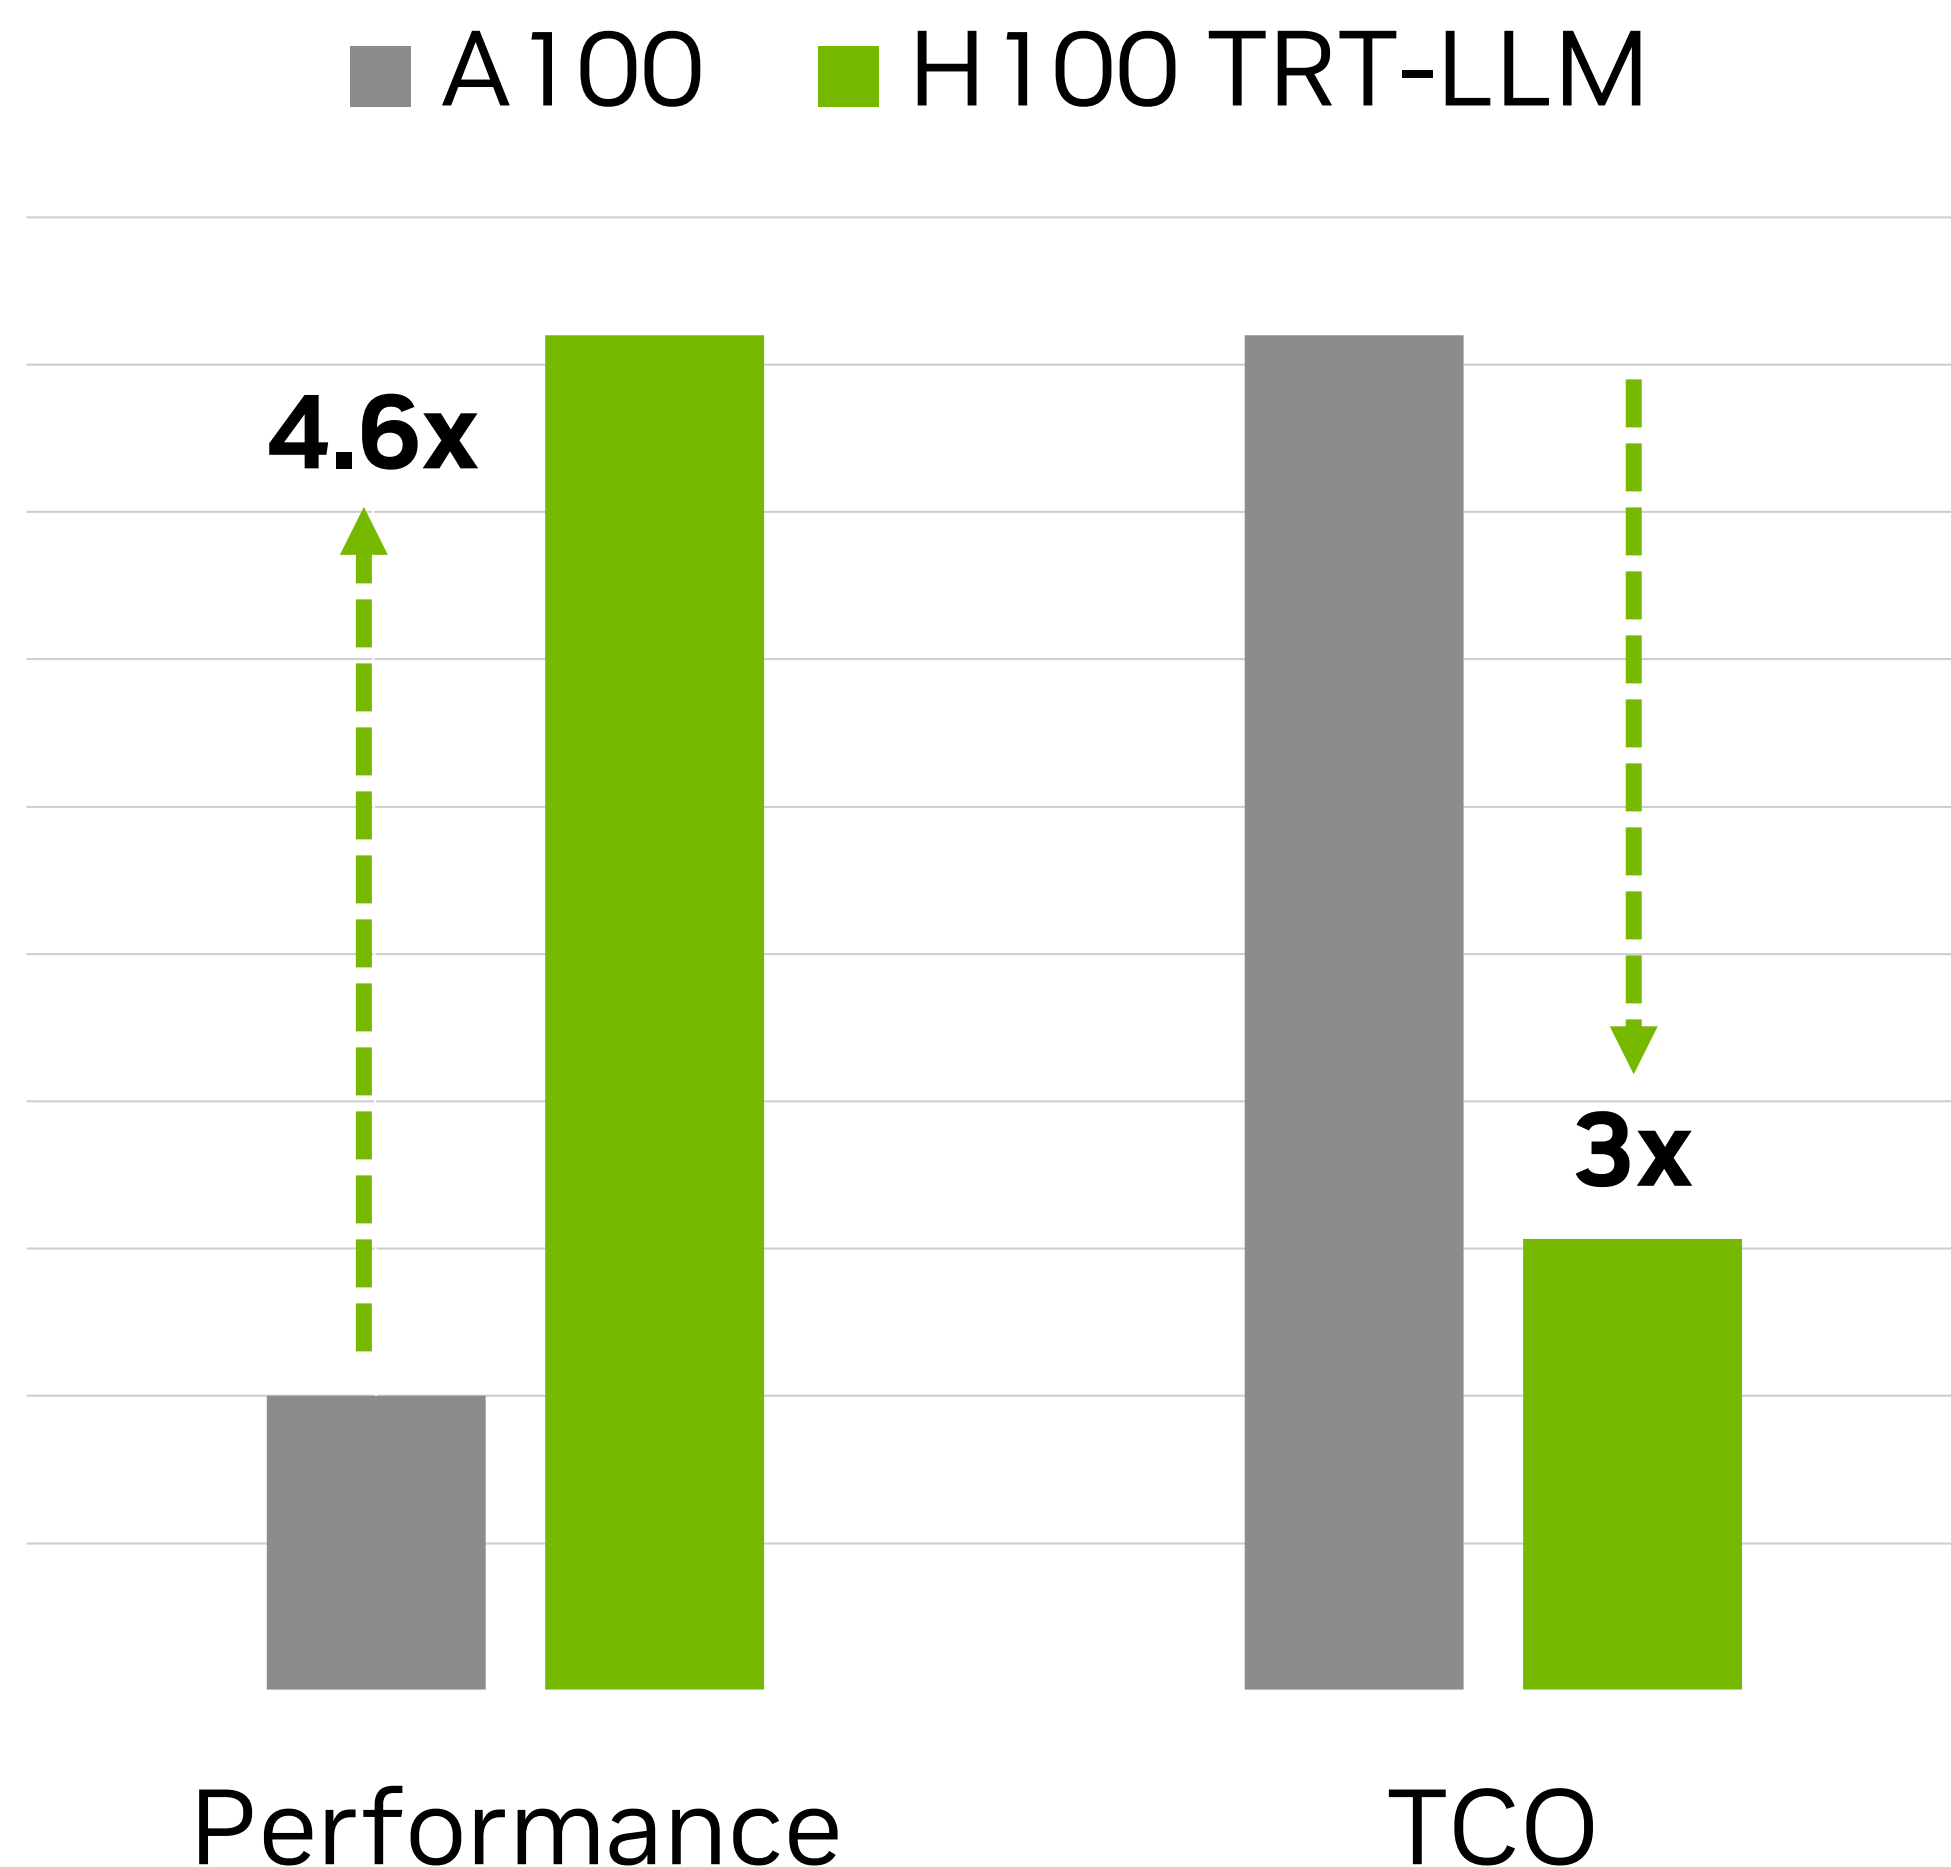
SoTA Performance for Large Language Models for Production Deployments

Challenges: LLM performance is crucial for real-time, cost-effective, production deployments. Rapid evolution in the LLM ecosystem, with new models & techniques released regularly, requires a performant, flexible solution to optimize models.

This presentation will focus on performance, and all the optimizations mentioned are available in TRT-LLM today!

SoTA Performance

Leverage TensorRT compilation & custom kernels



Ease Extension

Add new operators or models in Python to quickly support new LLMs with optimized performance

```
# define a new activation
def silu(input: Tensor) → Tensor:
    return input * sigmoid(input)

#implement models like in DL FWs
class LlamaModel(Module)
    def __init__(...)
        self.layers = ModuleList([...])

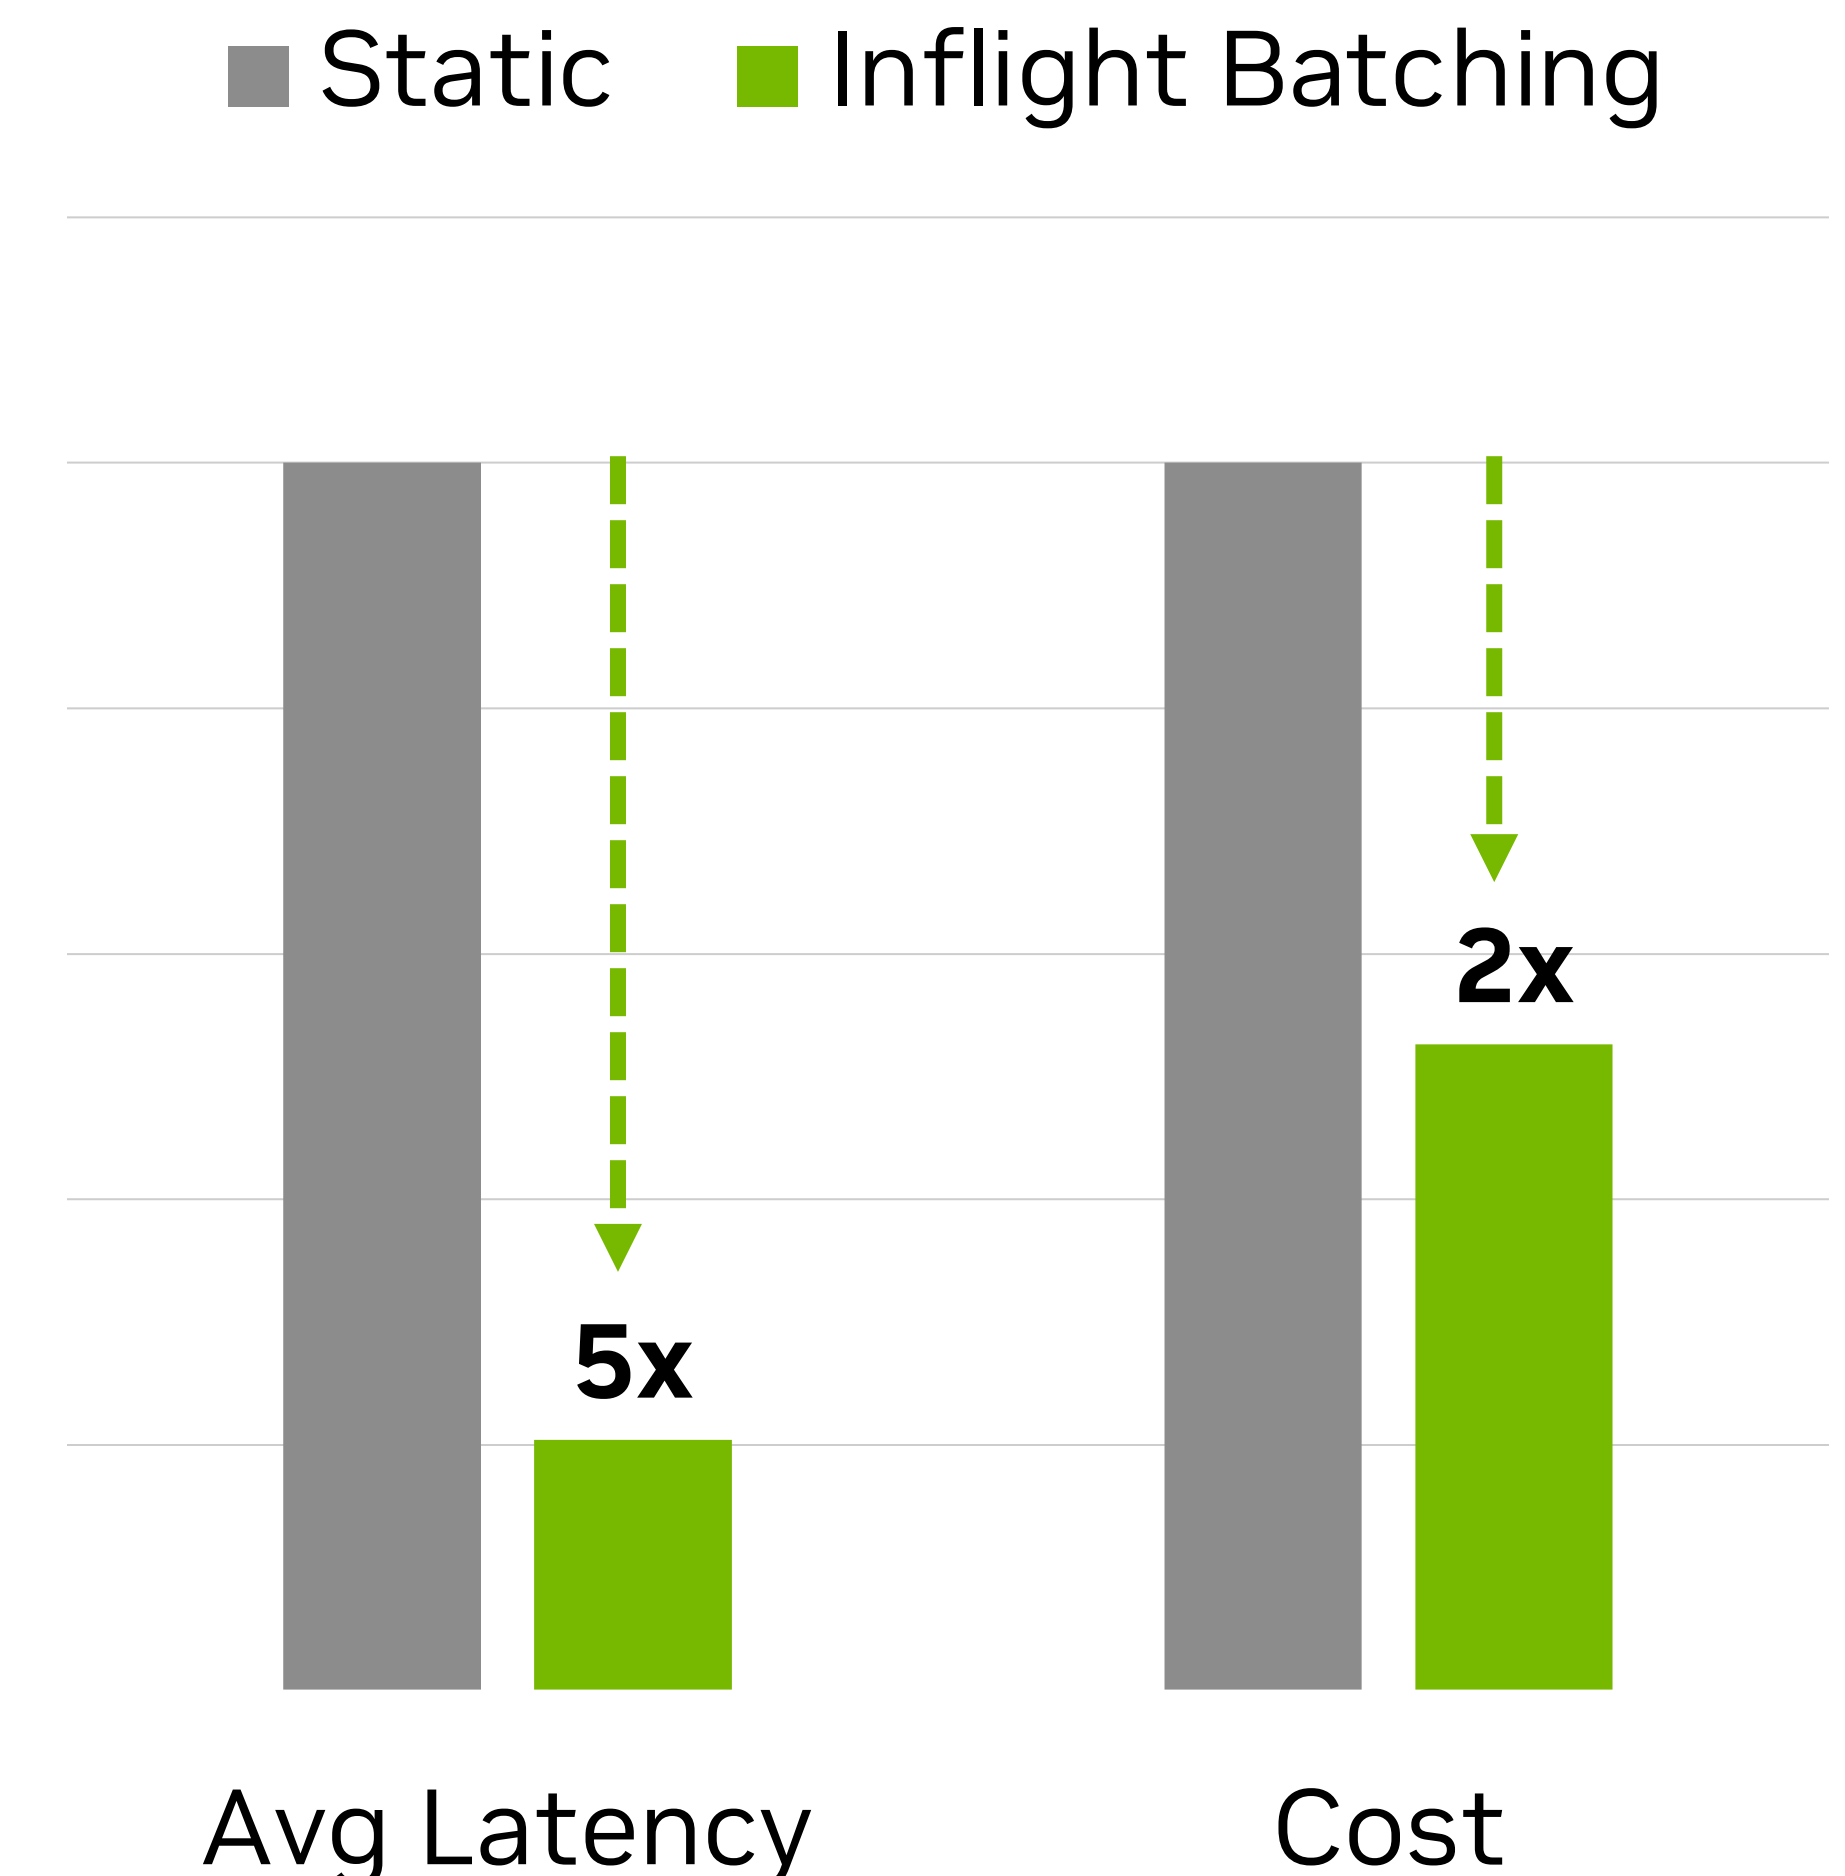
    def forward (...)
        hidden = self.embedding(...)

        for layer in self.layers:
            hidden_states = layer(hidden)

        return hidden
```

LLM Batching with Triton

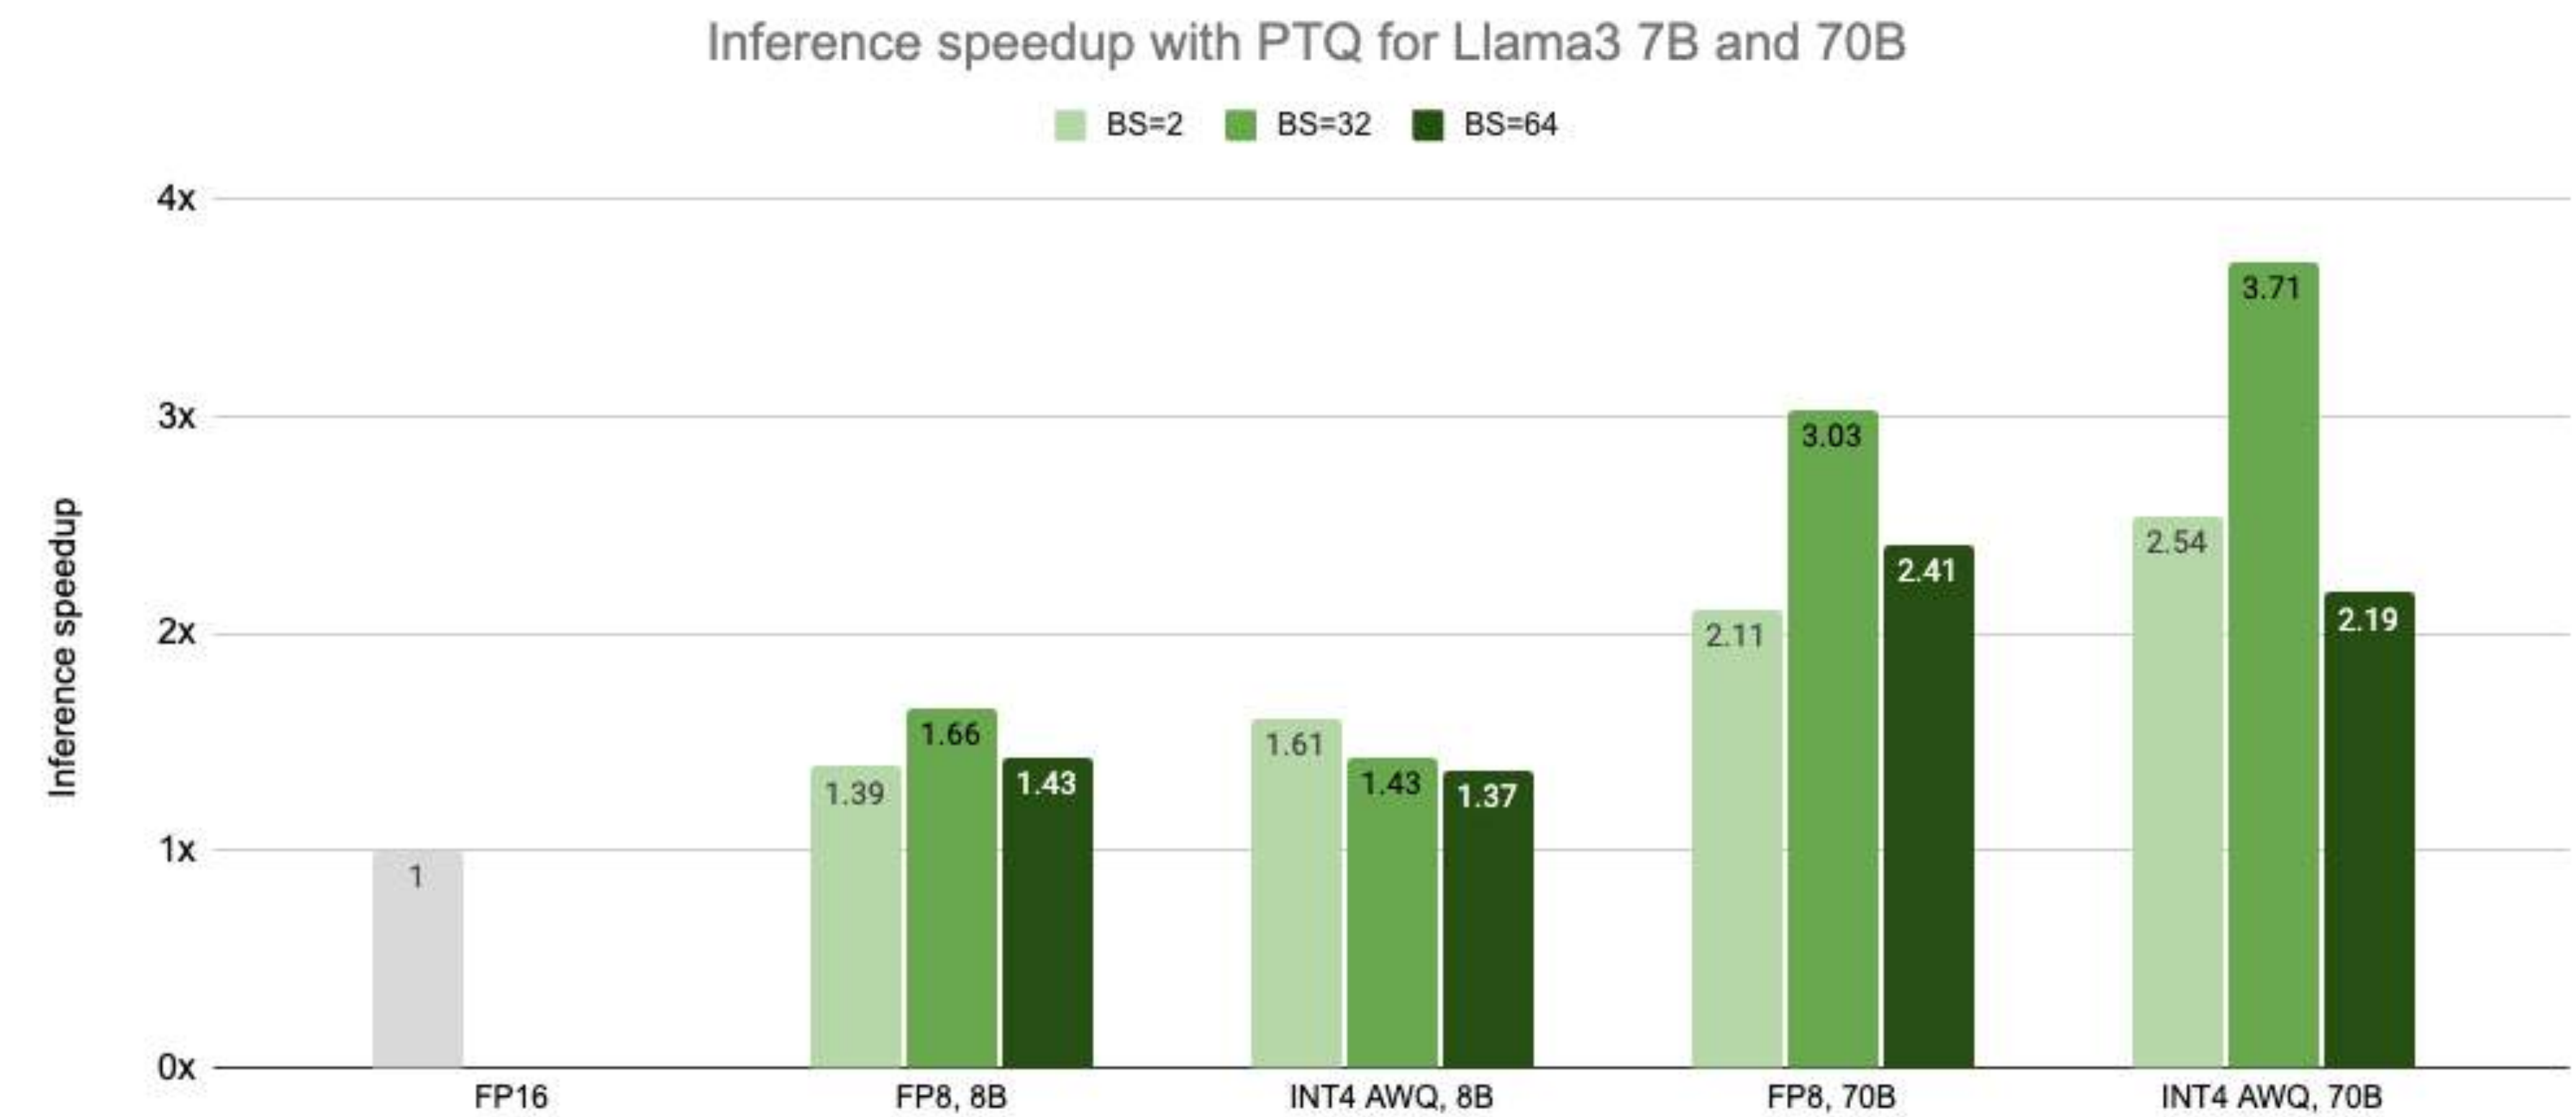
Maximize throughput and GPU utilization through new scheduling techniques for LLMs



Quantization

Lower precision representation

- BF16 is the default data precision in the world of LLMs
- Inference can be done in lower bit-width precisions - FP8, int8, int4, lower?
- Post-training quantization
- Result in lower DRAM BW pressure, lower memory footprint, higher throughput comms between GPUs and faster compute



TensorRT Model Optimizer Key Offerings

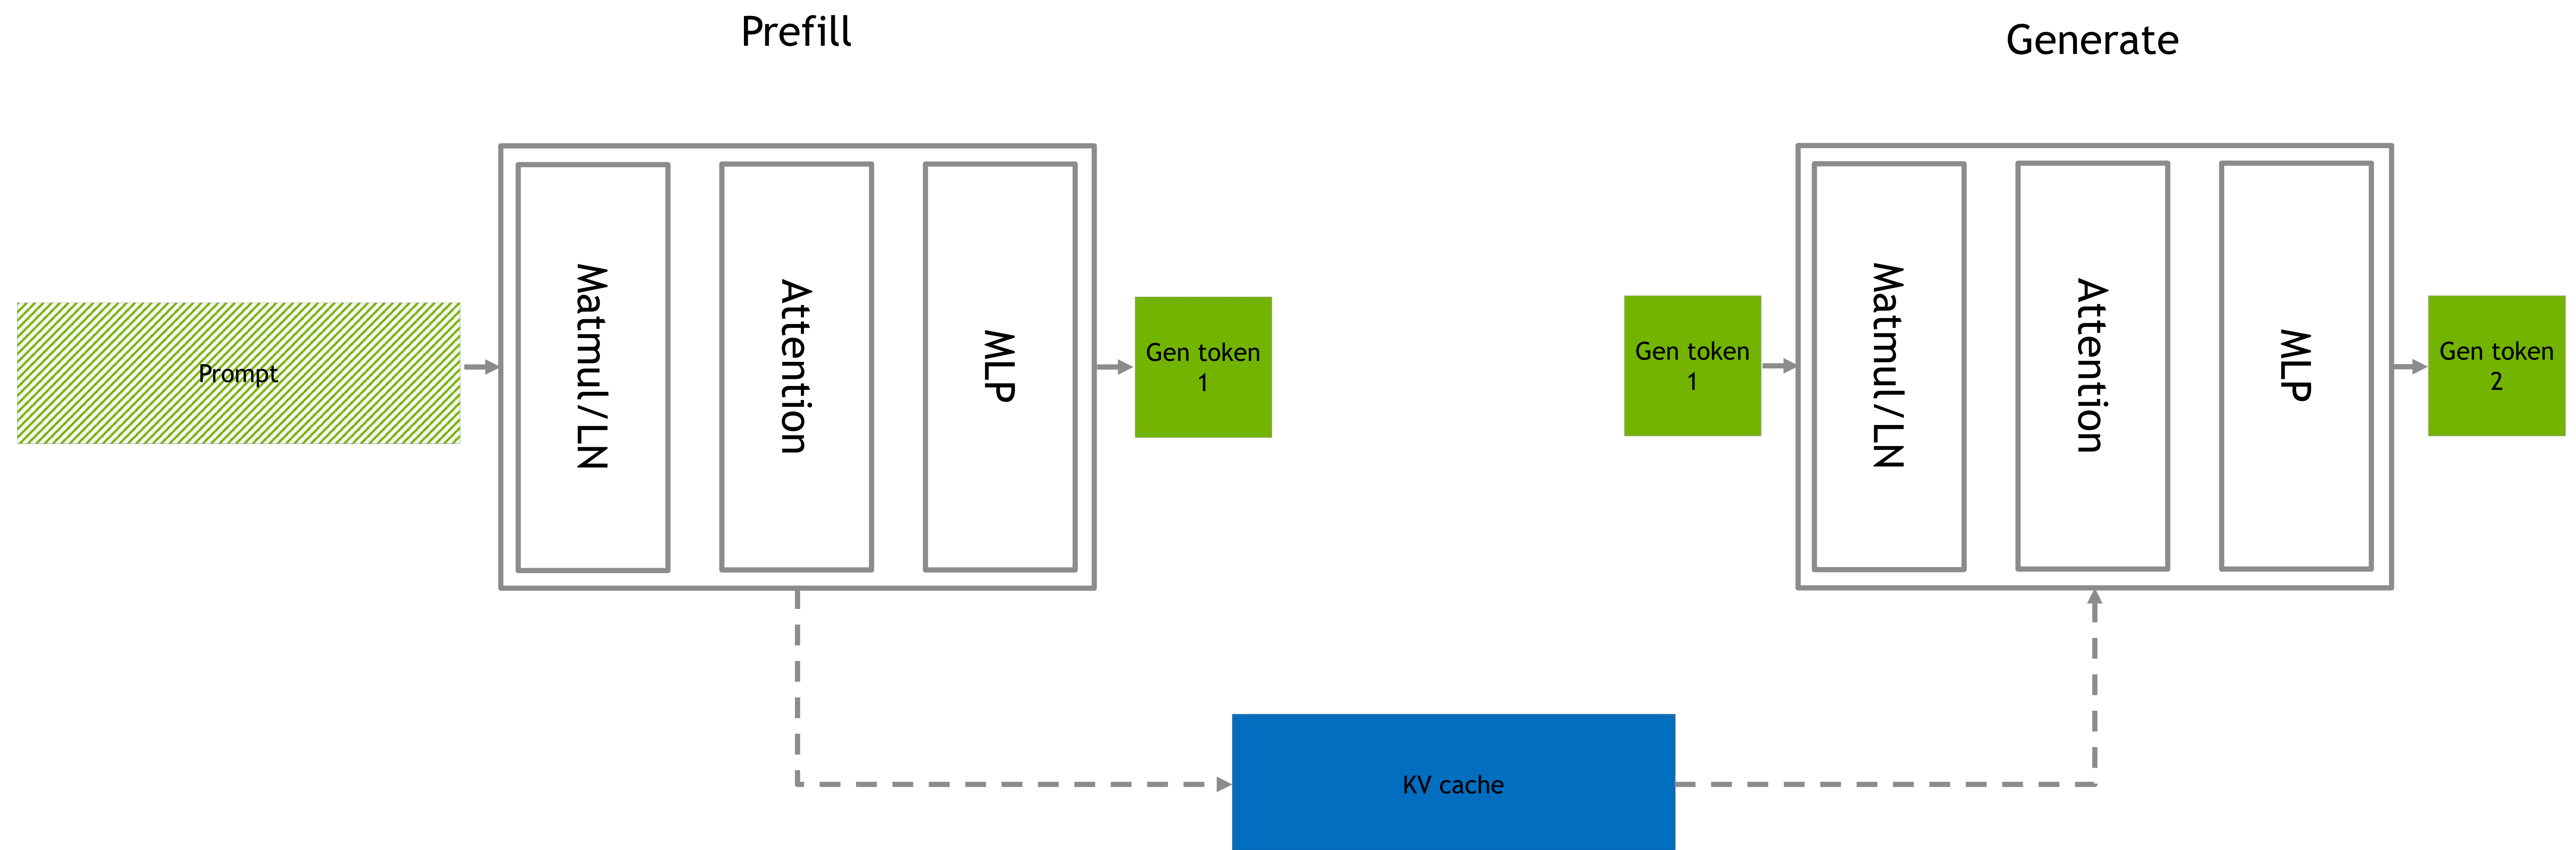
Latest release: [nvidia-modelopt v0.11](#)

- Techniques to optimize Generative AI models
 - Quantization
 - PTQ (Post-training quantization) and QAT (quantization aware training)
 - FP8, INT8, INT4. Supports advanced algorithms like SmoothQuant and INT4 AWQ
 - QAT examples with Hugging Face Trainer (New in v0.11 and soon w/MLM, NeMo)
 - Sparsity (New in v0.11)
 - Sparsity-aware fine tuning with Hugging Face Trainer
 - Post-training sparsity
- Easy-to-use & composable APIs
- Support PyTorch and ONNX models
- Native Windows support (New in v0.11)

Model	Quantization Methods	MMLU Baseline (FP16)	MMLU Post-quantization	MMLU Loss
Falcon-180B	FP8	70.4	70.3	0.14%
	INT8-SQ	70.4	68.6	2.56%
	INT4-AWQ	70.4	69.8	0.85%
Falcon-40B	FP8	56.1	55.6	0.89%
	INT8-SQ	56.1	54.7	2.50%
	INT4-AWQ	56.1	55.5	1.07%
LLaMA-v2-70B	FP8	69.1	68.5	0.87%
	INT8-SQ	69.1	67.2	2.75%
	INT4-AWQ	69.1	68.4	1.01%
MPT-30B	FP8	47.5	47.4	0.21%
	INT8-SQ	47.5	46.8	1.47%
	INT4-AWQ	47.5	46.5	2.11%

Phases of an LLM request

- Two phases – prefill and generate
- Prefill has lots of parallelism, and is commonly compute bound



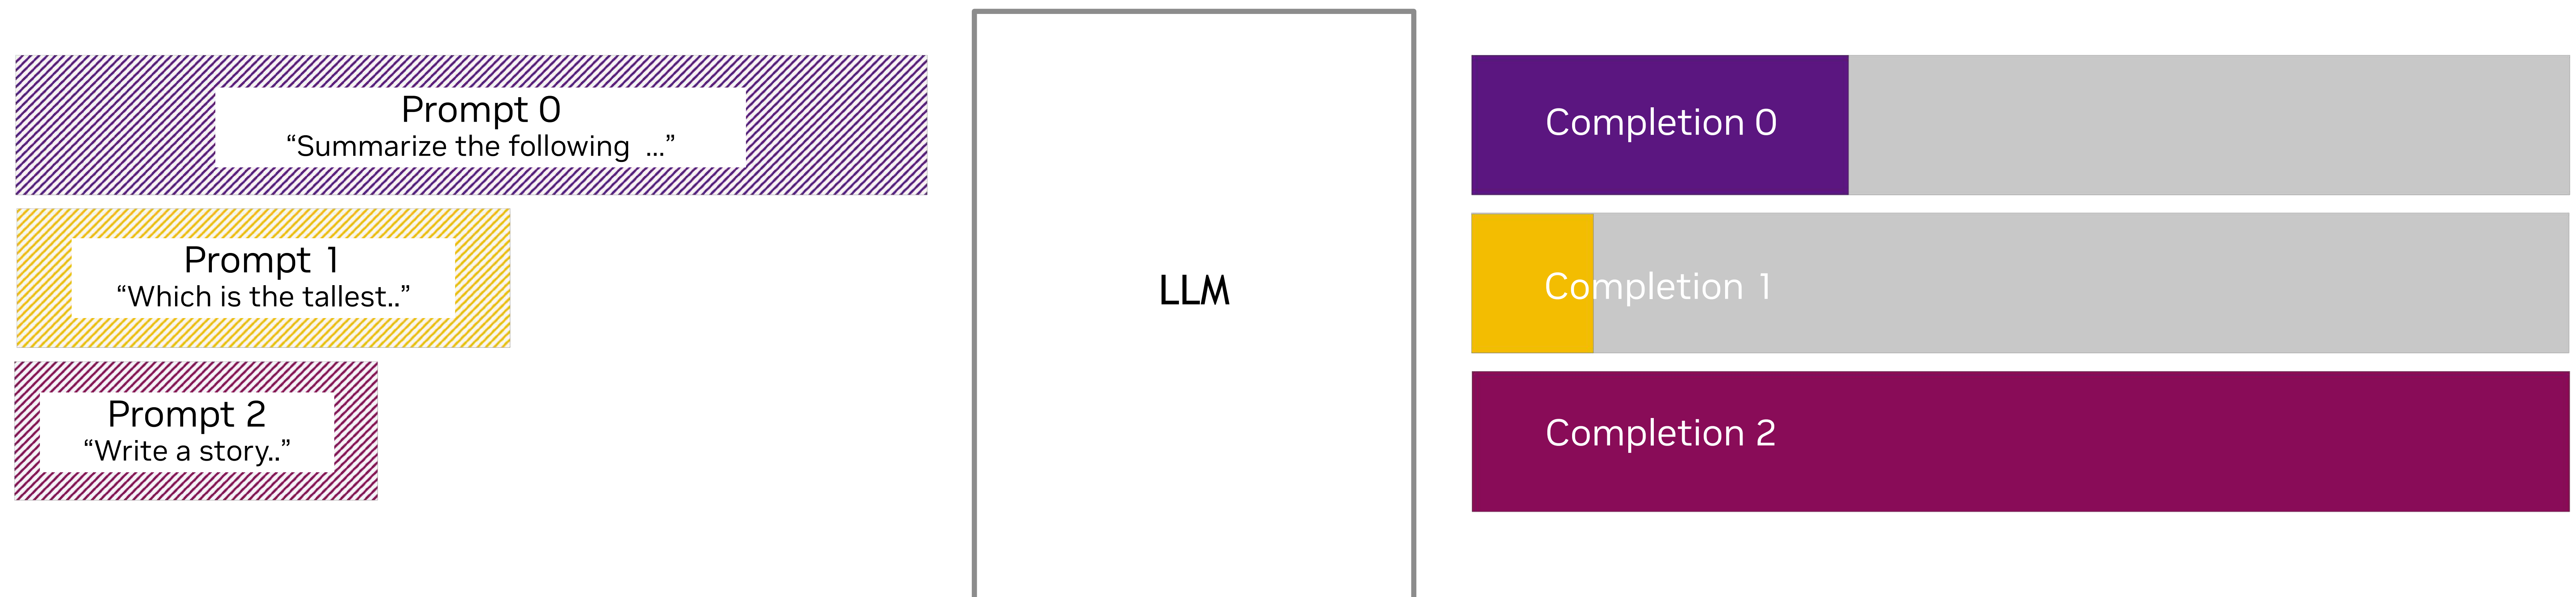
Optimized kernels

Fused MHA

Traditional request scheduling

Static batching

- Traditional view of Inference is that it is one forward prop of the network
- Requests are accumulated over some time window, and executed as a batch until completion
- “Batch size” view – some number of requests accepted, run in lockstep and are returned
- Is problematic when completions differ massively in length

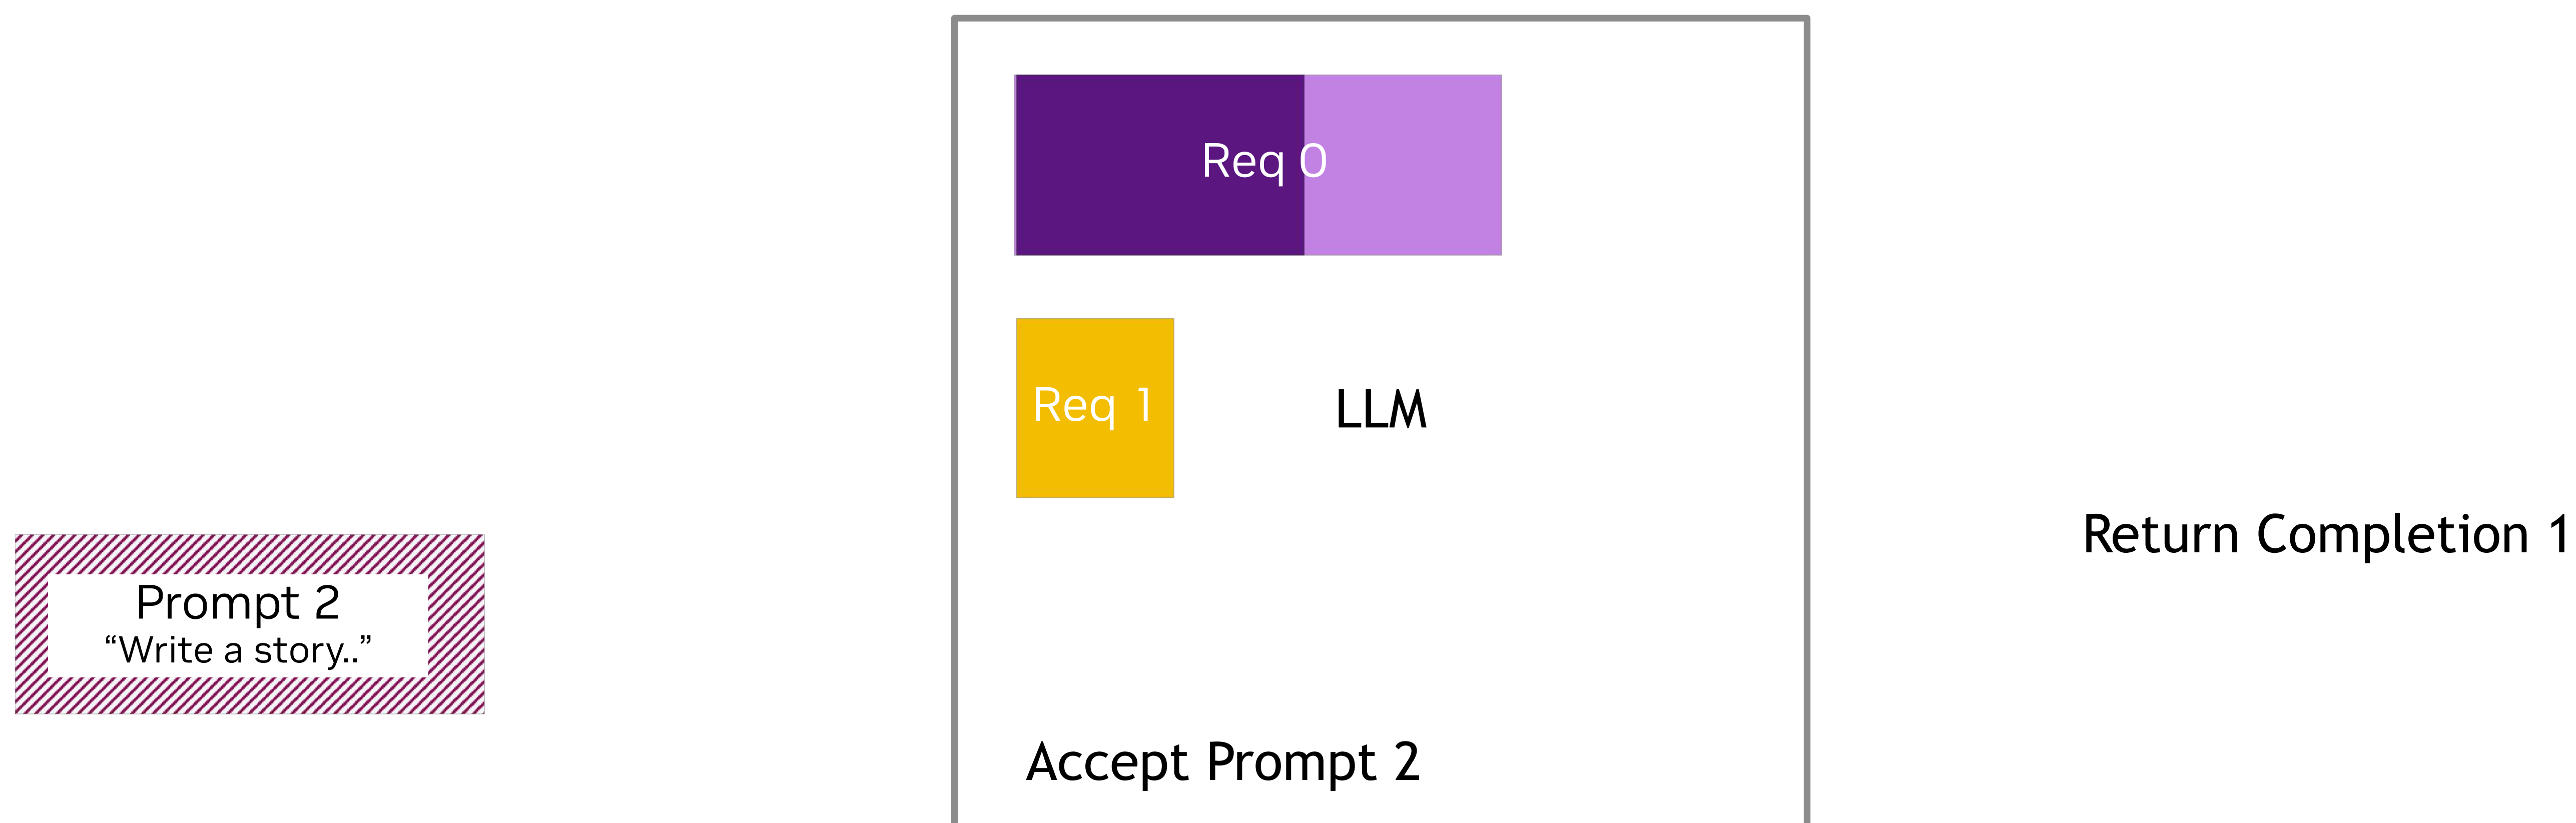


LLM-optimized request scheduling

In-flight batching

- Gen AI inference iteratively invokes multiple forward passes per request
- The number of passes is theoretically unbounded, and unknown a priori
 - You keep generating tokens until you produce and end signal
- Commonly in an online setting, so requests arrive over time, and latency matters
- LLM are Large, and expensive to run, so throughput matters to cost of operation
- You need a way to dynamically adapt to the workload instant-by-instant (or more correctly, iteration-by-iteration)

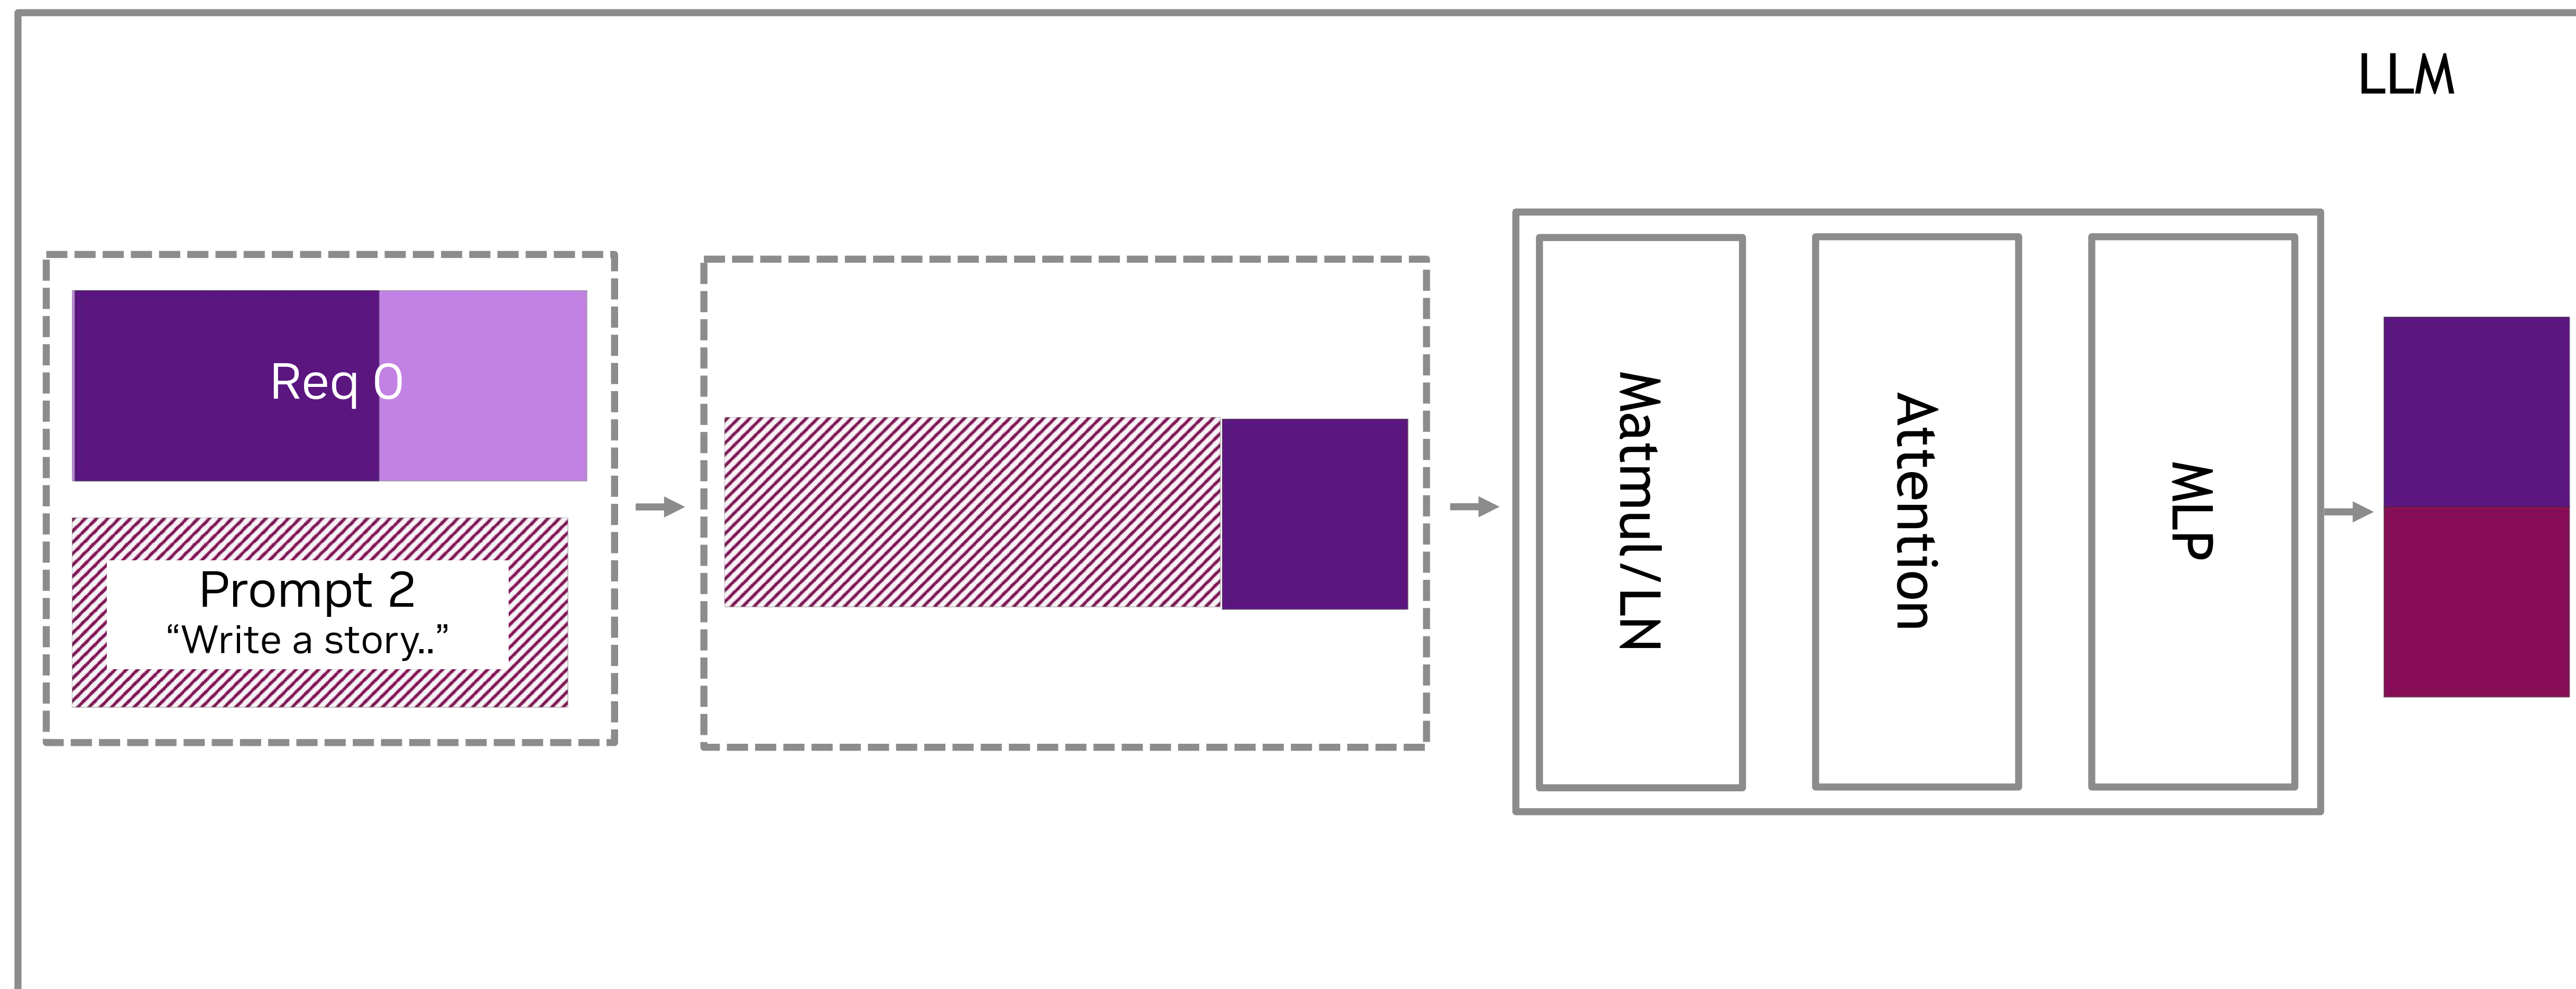
Snapshot at iteration i



Token concatenation

Batching affects the nature of the GPU workload

- Transformer models usually consist of two kinds of ops:
 - “token parallel” ops like Matmul and Layer Norm that operate just on individual tokens
 - “sequence parallel” ops like MHA that operate on sequences of data
- To a first order, LLM iteration runtime dominated by weights matrix multiplies, which are all token parallel
- Tokens across different requests in prefill and generate stages can be concatenated for higher matmul efficiency, and therefore, higher throughput

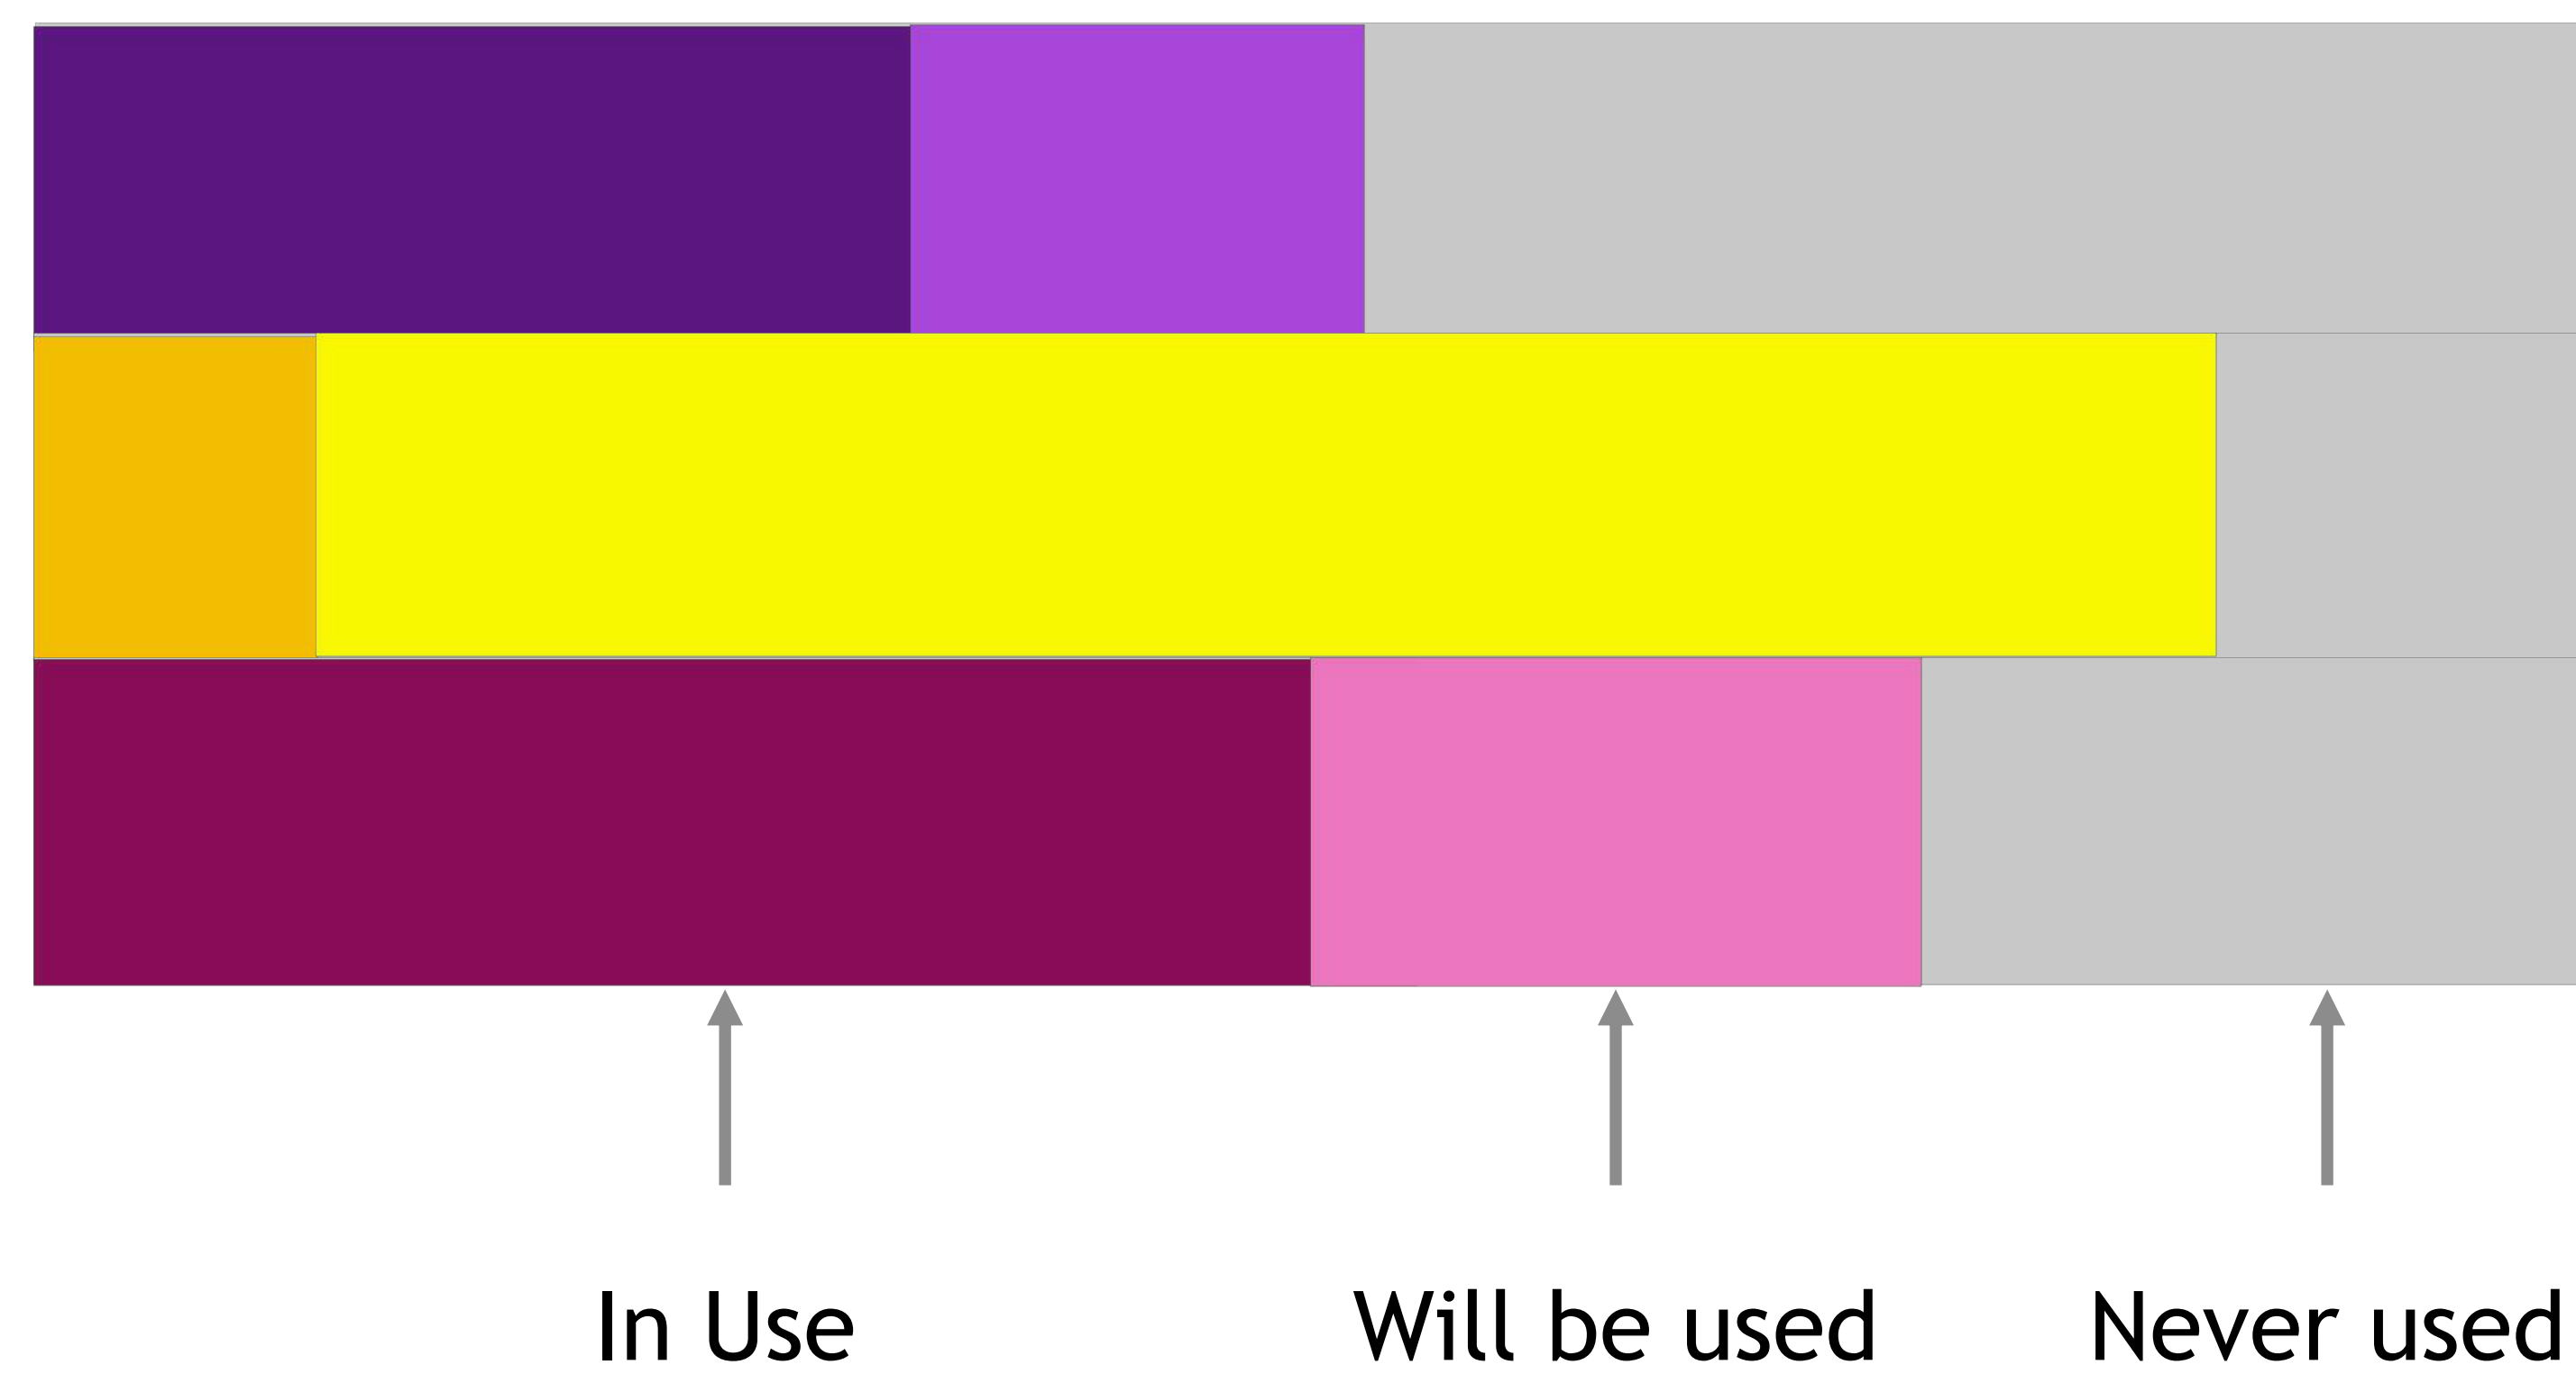


Paged KV cache

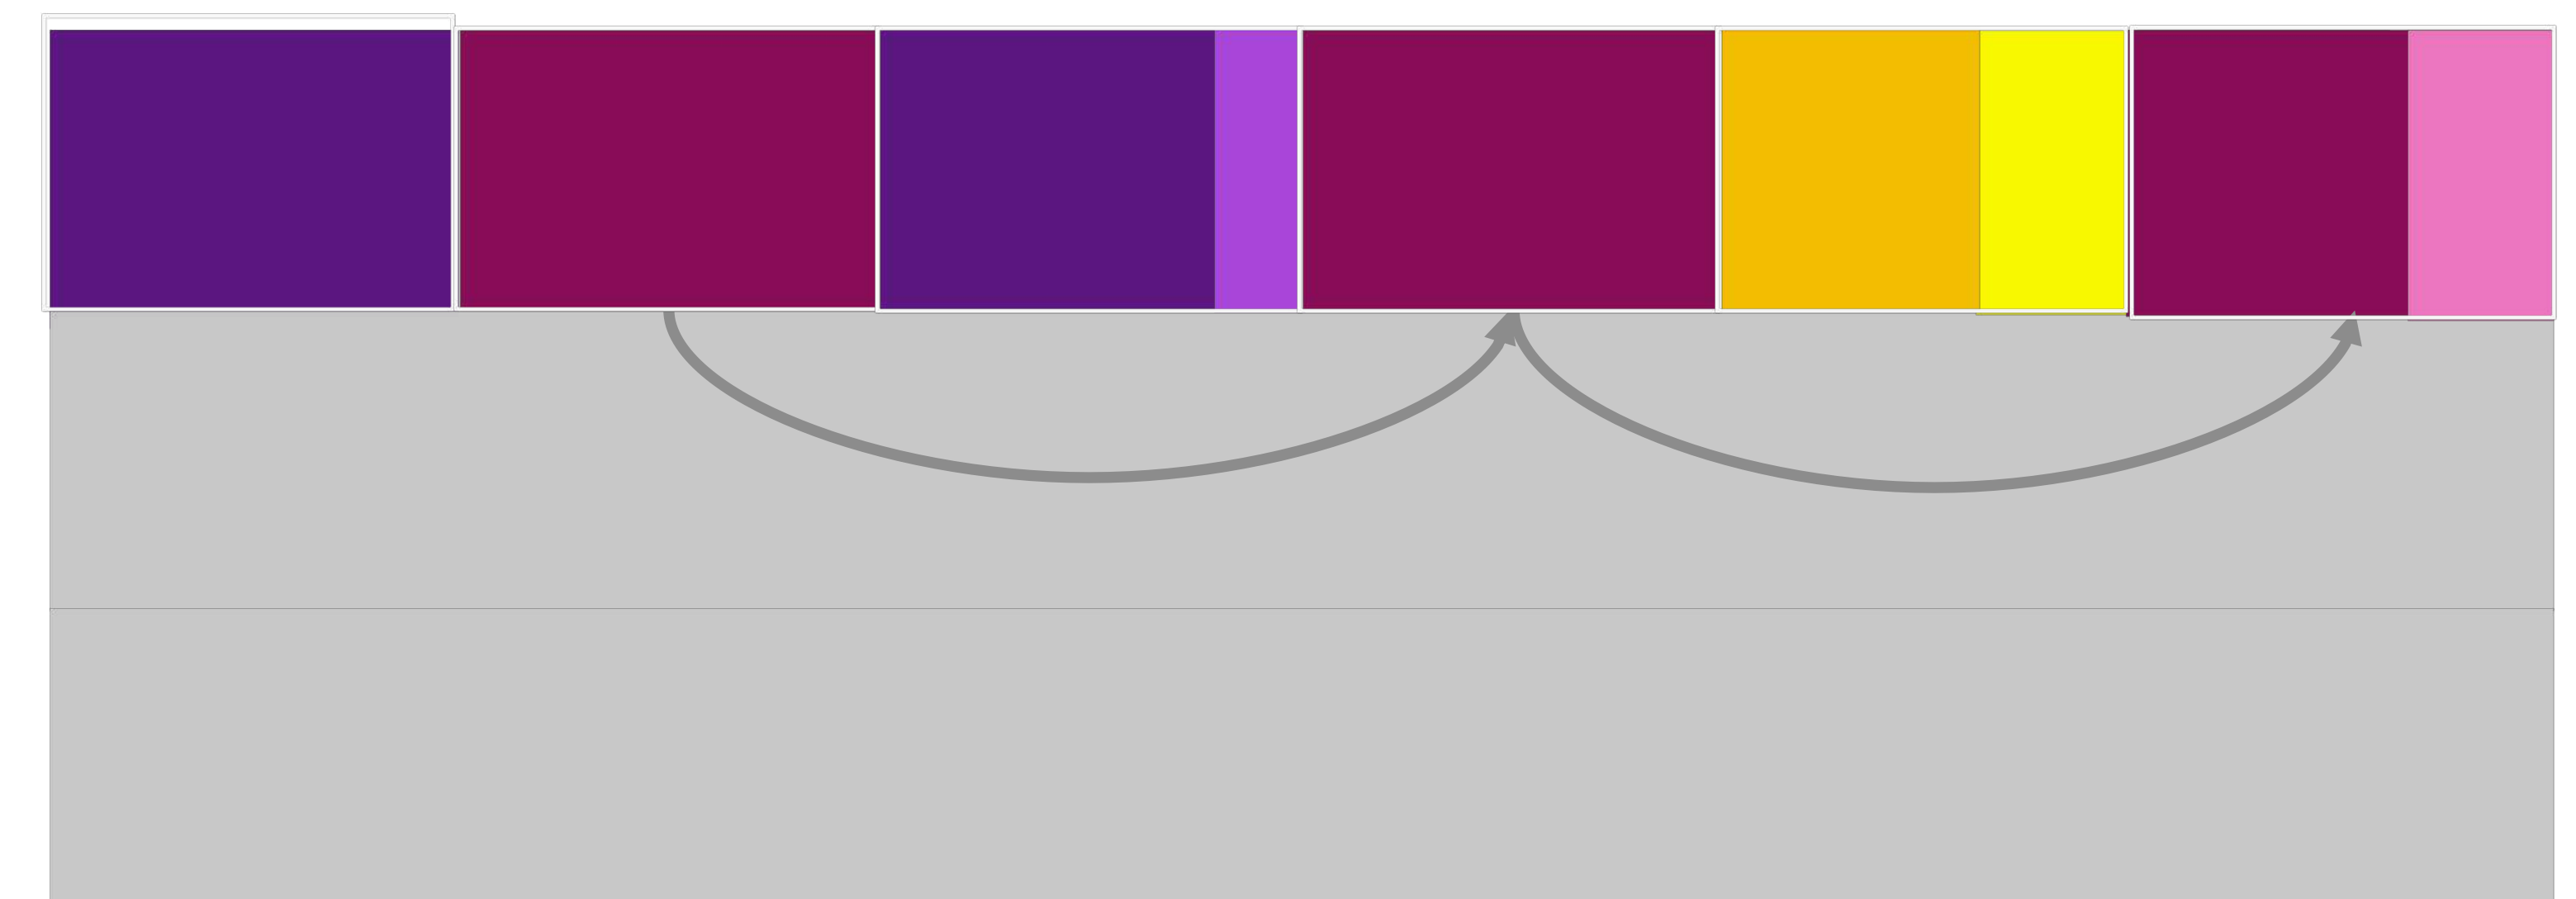
Optional subtitle

- Idea first proposed in vLLM
- Lazy memory allocation leads to minimal waste and more requests in-flight

Contiguous KV cache



Paged KV cache



KV cache reuse

Optional subtitle

- Linked list representation of paged KV cache allows for physical blocks to be shared across requests
- Cases like system prompts, where the first k tokens are shared across requests can benefit from memory and compute savings

