

A Crash Course on GPU Optimization – Mark Saroufim

Summary notes

1. Introduction:

- GPUs are expensive but valuable for accelerating machine learning models.
- The talk aims to provide techniques for getting the most out of GPUs (optimizing GPU utilization).
- Marc Saroufim is a PyTorch core developer at Meta and co-founder of the CUDA MODE Discord community.

2. PyTorch and GPU Acceleration:

- PyTorch is a tensor computation language with strong GPU acceleration.
- Using GPUs in PyTorch is as simple as calling ``.cuda()`` on tensors.
- Under the hood, PyTorch generates kernels for different devices, data types, and tensor layouts.
- The eager execution model in PyTorch (running line by line) can be less performant than a graph mode.

3. Pointwise Operations and CUDA Kernels:

- Pointwise operations (like ReLU) are common in deep learning models.
- PyTorch generates CUDA kernels for these operations, assigning threads to different memory locations.
- Thread scheduling strategies and taking advantage of the memory hierarchy are crucial for performance.

4. Techniques for GPU Optimization:

a. Fusing Operations:

- i. Instead of launching multiple small kernels, fusing operations into a single kernel can reduce overhead.
- ii. The ``.torch.compile`` function is a fusion compiler that generates fused Triton kernels.

b. Using Tensor Cores:

- i. Tensor Cores are special circuitry in NVIDIA GPUs for accelerating matrix operations.
- ii. Enabling Tensor Cores in PyTorch is done by setting ``.torch.set_float32_matmul_precision("high")``.

c. Reducing Overhead:

- i. Overhead can come from dispatching kernels and data transfers between CPU and GPU.
- ii. Queueing up multiple kernels in a CUDA graph using ``torch.compile(model, mode="reduce-overhead")`` can reduce overhead.

d. Quantization:

- i. Quantization can help both compute-bound and memory-bound workloads by reducing the data type size.
- ii. Weight-only quantization (int8 weights, bf16 activations) is a technique used in GPT-Fast.
- iii. The quantization space is rapidly evolving, with interest in sub-byte data types like int3, int4, etc.

e. Custom Kernels:

- i. For compute-bound problems, custom CUDA/Triton kernels may be required for optimal performance.
- ii. Techniques like Flash Attention and Online Softmax exploit GPU memory hierarchy and efficient normalization.
- iii. Learning CUDA and writing custom kernels can be rewarding but requires effort.

5. Learning Resources:

- The "Programming Massively Parallel Processors" (PMPP) book is recommended for learning CUDA.
- The ``torch.utils.cpp_extension.load_inline()`` function simplifies writing and testing CUDA kernels within PyTorch.
- The NVIDIA Nsight Compute profiler (ncu) is a valuable tool for analyzing kernel performance.
- The CUDA MODE Discord community and lectures are great resources for learning and collaboration.

6. Concluding Thoughts:

- Performance optimization is important as demand for faster models continues to grow.
- While compilers excel at fusions, kernel authors are still needed for rewriting mathematical operations efficiently.
- Marc is personally focused on quantization and custom kernels in the AO repo and CUDA MODE community.