




VOLTRON DATA

Scaling data processing from CPU to distributed GPUs

By William Malpica,
Director of Engineering and
Co-Founder at Voltron Data



• CPUs vs GPUs

CPUs are great at everything!

- ... well good enough at most things
- Generic, easy to use, found everywhere and less expensive than GPUs
- Fast and low latency
- Efficient at doing lots of different things
- They are like USPS



Generic Serial Processing



GPUs are even better

- ... at specific things....
- GPUs excel at doing the same thing over and over in parallel
- Prefer data to be very structured in vectors and arrays
- High latency, high throughput
- They are like shipping trucking company



Massively Parallel Processing



• Data processing... what kind of data processing?

OLTP (Online Transactional Processing)

- Row databases are commonly used for OLTP, where a single “transaction,” such as inserting, removing or updating, can be performed quickly using small amounts of data.
- Each row can hold very different datatypes and each row can be different sizes in memory
- Classic Row databases include MySQL and PostgreSQL
- Classic Row-wise file formats include: CSV, TFRecord file
- Not GPU friendly

A Row of packages?



OLAP (Online Analytical Processing)

- Column based databases are commonly used for OLAP, where each column is held in contiguous memory
- Column based databases make it awkward to insert or remove data, but it makes it much easier to do analytics.
- Classic Column databases include Snowflake and AWS Redshift.
- Classic Column-wise file formats include: Apache Parquet and Apache ORC
- SIMD and GPU friendly


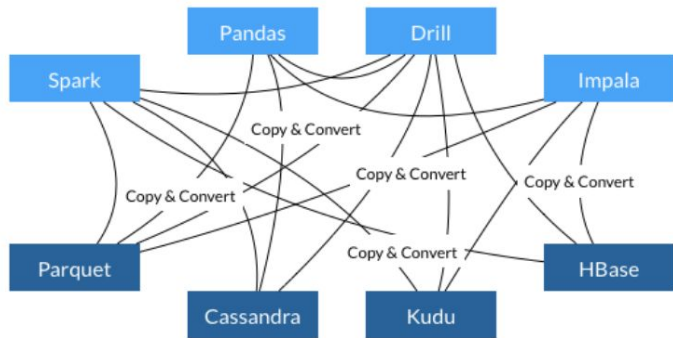
Columns of packages?



• Meet Apache Arrow

- It's a columnar memory format
- Its language-independent
- Designed for efficient analytic operations on modern CPUs and GPUs

- Supports zero-copy data transfers
- Its the de-facto standard of memory representation in modern data analytics
- Over 100M downloads a month



**GPUs can
accelerate many
types of Data
Processing**

Single GPU DataFrame operations: cuDF vs. Pandas


cuDF v0.10, Pandas 0.24.2

□ Running on NVIDIA DGX-1:

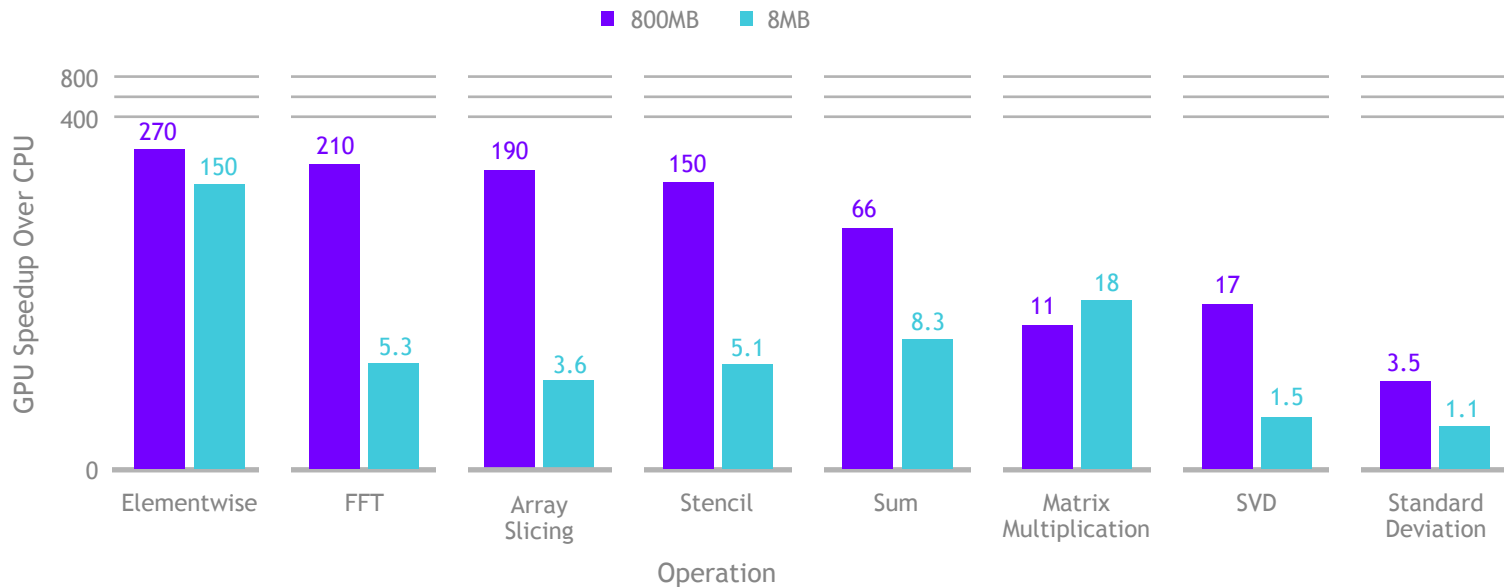
- GPU: NVIDIA Tesla V100 32GB
- CPU: Intel(R) Xeon(R) CPU E5-2698 v4 @ 2.20GHz

□ Benchmark Setup:

- DataFrames: 2x int32 columns key columns, 3x int32 value columns
- Merge: Inner
- GroupBy: Count, sum, min, max calculated for each value column

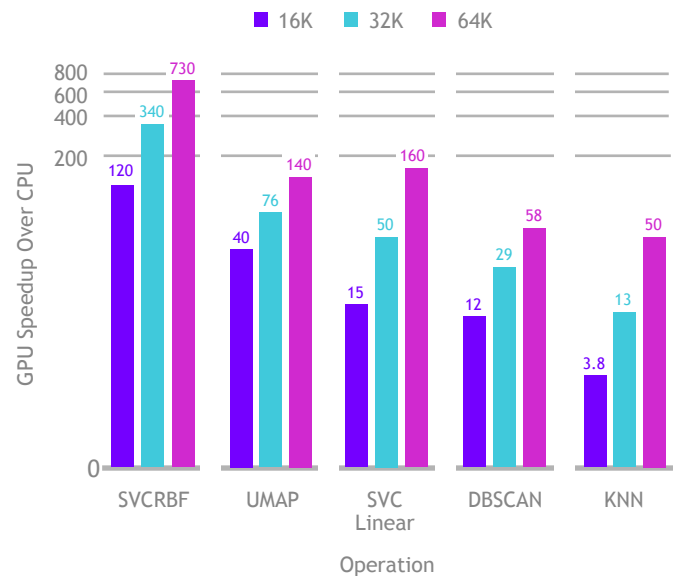



Single GPU Numeric Array operations: CuPy vs NumPy



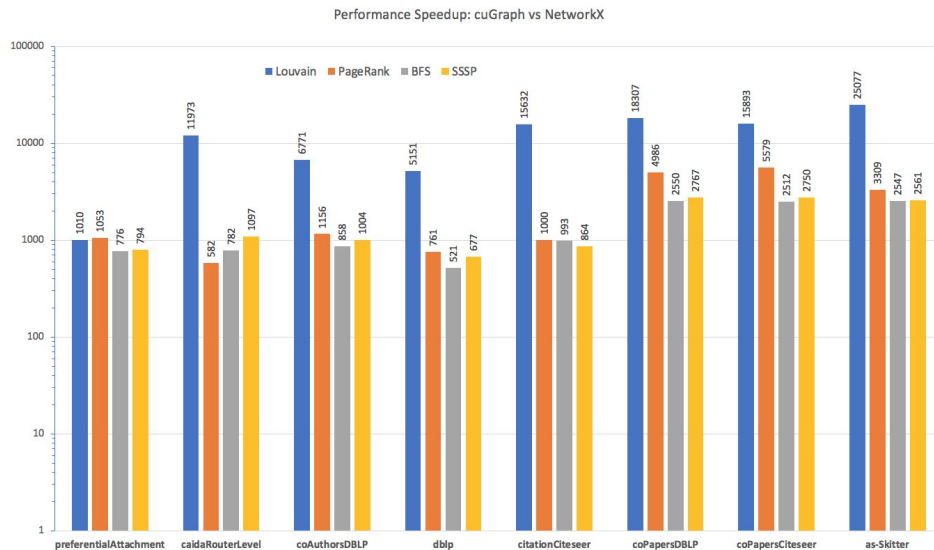
More details: <https://blog.dask.org/2019/06/27/single-gpu-cupy-benchmarks>

Single GPU Machine Learning: cuML vs Scikit-learn



1x V100 vs. 2x 20 Core CPU

Single GPU Graph Analytics: cuGraph vs NetworkX




Dataset	Nodes	Edges
preferentialAttachment	100,000	999,970
caidaRouterLevel	192,244	1,218,132
coAuthorsDBLP	299,067	299,067
DBlp-2010	326,186	1,615,400
citationCiteseer	268,495	2,313,294
coPapersDBLP	540,486	20,491,458
coPapersCiteseer	434,102	32,073,440
As-Skitter	1,696,415	22,190,596

- **Other GPU accelerated libraries aside from RAPIDS:**



- ▶ **pyTorch**
 - ♦ Deep learning
- ▶ **ArrayFire**
 - ♦ General parallel computing
- ▶ **TensorFlow**
 - ♦ Machine learning platform
- ▶ **Numba**
 - ♦ Is a just-in-time compiler for accelerating python code
- ▶ **Bend**
 - ♦ A massively parallel, high-level programming language.
- ▶ **XGBoost**
 - ♦ ML library which uses Distributed gradient boosting
- ▶ **Voltron Data Theseus SQL Engine**
 - ♦ SQL data analytics

How are GPUs so fast?

- # Number of Cores



- # Memory Bandwidth



Ok, so GPUs are fast.

Should I always use them?

• When GPUs are a good idea?

CPUs might be a better idea:

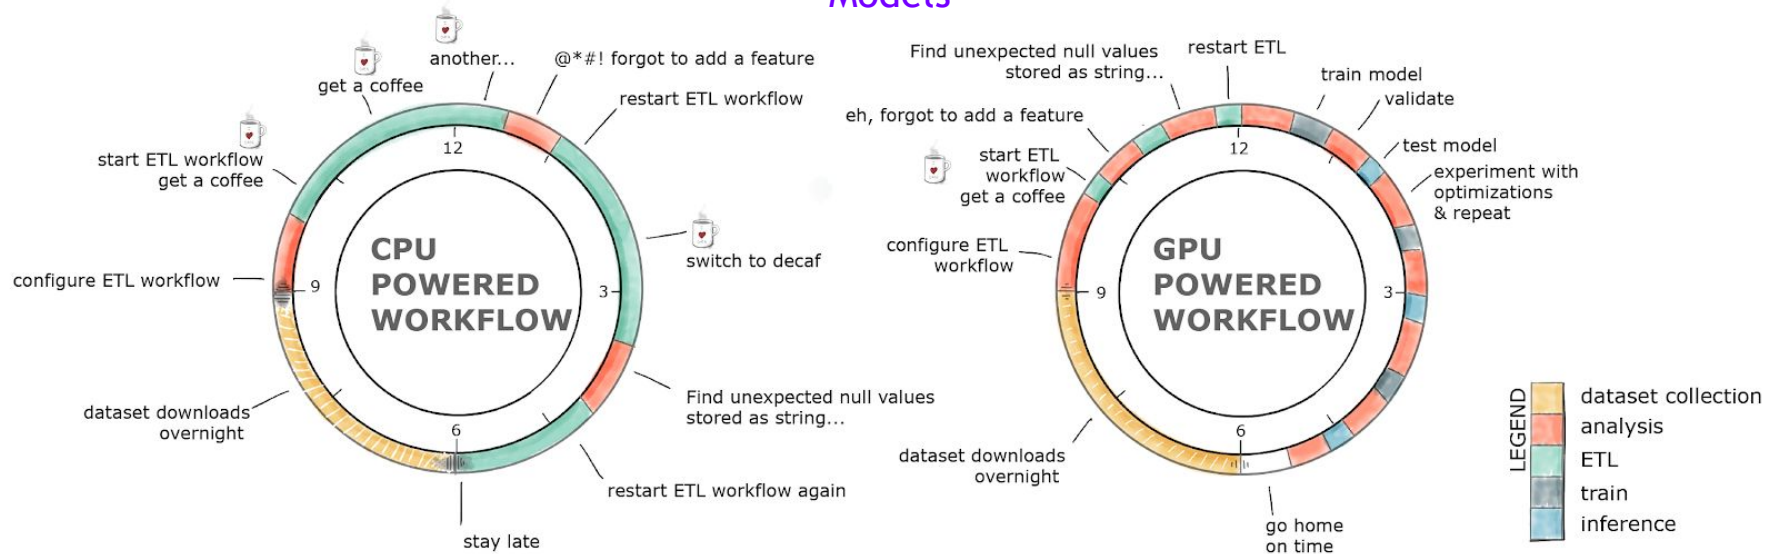
- If the data processing is latency bound
- If the data processing is I/O bound
- If the costs using GPUs outweigh the benefits
- If the amount of data is not too much

GPUs might be a better idea if:

- If your data processing is compute bound
 - ▶ NOTE: When using GPUs you may become I/O bound. It can be hard to feed the beast!
- If compute density is important to you (more on this later)
- If throughput speed enables you to do more
- If you are using GPUs anyways

GPU Powered workflow can reduce iteration time

The Average Data Scientist Spends 90+% of Their Time in ETL as Opposed to Training Models




**Large data
problems usually
require distributed
solutions**

• The Wall

CPU performance is capped. No amount of money will jump over this wall.

TPC-H 10TB Benchmark

● Spark EMR



Note: Theseus: 1 Node 8 x A100 80 GB, Spark: 1 Node r5.8xlarge (AWS) 32 vCPU 32 GB

GPUs can help jump over the wall

TPC-H (10TB, 30TB, 100TB)

- ✓ Up to 10 DGX Servers
- ✓ Parquet Files
- ✓ Remote File System
- ✓ Lots of Spilling


- ✗ No Sorting
- ✗ No Indexing
- ✗ No Caching
- ✗ No Warm Up (Cold Queries)

Note: Theseus: 1 Node 8 x A100 80 GB,
Spark: 1 Node r5.8xlarge (AWS) 32 VCPU
32 GB

SPACE: Scale, performance, and cost efficiency


Theseus 10TB Theseus 30TB Theseus 100TB Spark 10TB Spark 30TB

↑ Total Runtime (minutes)




• New Bottleneck

Speeding up the compute just moves the bottleneck elsewhere, Networking and Storage



- # Networking and Storage

Need to be aware of the hardware when architecting the software




• How to get the most performance on your distributed GPU data processing system?

- Use GPUDirect RDMA to communicate with Infiniband or RoCE
- Use GPUDirect Storage to write to local or remote storage, such as NVMe or NVMe over Fabric (NVMe-oF)
- It avoids extra copies through a bounce buffer in the CPU's memory, enabling a direct memory access (DMA) to move data on a direct path into or out of GPU memory — all without burdening the CPU.
- Need to use distributed systems which can support these types of technologies:
 - ▶ Dask: can be configured for GPUDirect RDMA with OpenUCX
 - ▶ Spark Rapids: can use GPUDirect RDMS and GPUDirect Storage
 - ▶ MPI: can be configured for GPUDirect RDMA
 - ▶ Theseus: will use GPUDirect RDMS and GPUDirect Storage when available

Voltron Data Theseus

A Compute Mesh unifying hardware, languages, and applications



- 1 Accelerator-Native:**
Distributed query engine built from the ground up to take advantage of full system hardware acceleration.
- 2 Petabyte Scale:**
Focusing on problems too big and time sensitive for Spark
- 3 Composable:**
Built on open source standards that enables interoperability from storage to application
- 4 Evolutionary:**
A composable engine that seamlessly adapts to new hardware and languages

A low-angle, upward-looking photograph of several modern skyscrapers. The image is characterized by strong geometric lines and patterns, including a series of parallel diagonal lines from a building's facade and a grid-like pattern on another. Two teal-colored geometric shapes, a circle and a triangle, are overlaid on the image. The text "Thank You!" is written in white on the teal circle, and "Questions?" is written in white on the teal triangle.

Thank You!

Questions?